

# IMPERIAL

IMPERIAL COLLEGE LONDON

DEPARTMENT OF MATHEMATICS

MSCI RESEARCH PROJECT

---

## **A Variational Approach to American Option Pricing via Physics-Informed Neural Networks**

---

*Author:*  
James Wu

*Supervisor(s):*  
Harry Zheng

Submitted in partial fulfillment of the requirements for the MSci in Mathematics at Imperial College London

June 6, 2026

## **Abstract**

American option pricing requires solving free boundary PDEs with no closed-form solution. We reformulate the problem as a variational inequality and propose a physics-informed neural network (PINN) approach to solve it. On both the Black-Scholes and Heston models, our PINNs match tree-based benchmarks to within 3% relative error. We extend the approach to multi-asset settings where we benchmark against equivalent one-dimensional reductions. We characterise the free boundary from the learned solution, yielding interpretable early-exercise surfaces. The results establish PINNs as a practical alternative to lattice methods for higher-dimensional problems in financial mathematics.

## Acknowledgments

I would like to express my sincere gratitude to my supervisor, Professor Harry Zheng, for his expert guidance and continued support throughout this project. Our weekly meetings and his insights were invaluable in the development of this research.

I am also grateful to Yun Zhao for his continued assistance, especially in the numerical implementation side of the project. His help was crucial in overcoming the technical challenges we faced.

Finally, I would like to thank my family for their patience and encouragement throughout what has been a challenging degree. They have provided a constant source of motivation, which I deeply appreciate.

## Plagiarism statement and use of AI

The work contained in this thesis is my own work unless otherwise stated.

The use of AI tools in this dissertation is limited. I used ChatGPT and Claude to help me understand some of the mathematical concepts and methods used, but my research was done manually through published, peer-reviewed sources, all of which are cited in the bibliography. My implementation of the methods in Python are initially done without AI, but are later debugged and optimised with the help of Claude Code.

*Signature:* James Wu

*Date:* June 6, 2026

# Contents

|          |                                                       |           |
|----------|-------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                                   | <b>8</b>  |
| 1.1      | Motivation . . . . .                                  | 8         |
| 1.2      | Literature review . . . . .                           | 9         |
| 1.3      | Overview of project . . . . .                         | 10        |
| <b>2</b> | <b>Mathematical Background</b>                        | <b>12</b> |
| 2.1      | Preliminaries . . . . .                               | 12        |
| 2.1.1    | Probabilistic Setup . . . . .                         | 12        |
| 2.1.2    | Stochastic processes . . . . .                        | 13        |
| 2.1.3    | Risk-neutral pricing . . . . .                        | 14        |
| 2.1.4    | Optimal stopping problem and free boundary . . . . .  | 14        |
| 2.1.5    | Monte Carlo integration . . . . .                     | 17        |
| 2.2      | State Space Discretisation . . . . .                  | 18        |
| 2.2.1    | Background . . . . .                                  | 18        |
| 2.2.2    | Tree Methods . . . . .                                | 18        |
| 2.3      | Physics-Informed Neural Networks . . . . .            | 19        |
| 2.3.1    | Motivation . . . . .                                  | 19        |
| 2.3.2    | Methodology . . . . .                                 | 20        |
| 2.3.3    | Adaptive Loss Weighting . . . . .                     | 21        |
| 2.3.4    | Numerical Optimisation . . . . .                      | 22        |
| <b>3</b> | <b>Black-Scholes Model</b>                            | <b>24</b> |
| 3.1      | Closed-form solution for European options . . . . .   | 24        |
| 3.2      | Binomial Tree Method . . . . .                        | 25        |
| 3.3      | Physics-Informed Neural Networks . . . . .            | 27        |
| 3.3.1    | PINN formulation for one-dimensional case . . . . .   | 27        |
| 3.3.2    | Multi-asset case . . . . .                            | 29        |
| 3.3.3    | Reduction of multi-asset to one-asset case . . . . .  | 29        |
| 3.3.4    | PINN formulation for multi-dimensional case . . . . . | 31        |

|          |                                                       |           |
|----------|-------------------------------------------------------|-----------|
| 3.4      | Sobolev Deep Neural Networks . . . . .                | 32        |
| 3.4.1    | Sobolev spaces . . . . .                              | 33        |
| 3.4.2    | Model . . . . .                                       | 33        |
| 3.4.3    | Loss function . . . . .                               | 34        |
| 3.5      | Numerical Results . . . . .                           | 35        |
| 3.5.1    | Tree benchmark against closed-form solution . . . . . | 35        |
| 3.5.2    | PINN results against tree benchmark . . . . .         | 36        |
| 3.5.3    | Multi-dimensional PINN . . . . .                      | 38        |
| 3.5.4    | Sobolev PINN results . . . . .                        | 40        |
| 3.5.5    | Free boundary . . . . .                               | 42        |
| 3.6      | Minimum of two assets . . . . .                       | 44        |
| <b>4</b> | <b>The Heston Model</b>                               | <b>47</b> |
| 4.1      | Closed-form Solution . . . . .                        | 47        |
| 4.2      | Tree Solution . . . . .                               | 49        |
| 4.2.1    | Euler scheme . . . . .                                | 50        |
| 4.2.2    | Interpolation optimisation . . . . .                  | 50        |
| 4.2.3    | Backwards induction . . . . .                         | 52        |
| 4.2.4    | Implementation techniques . . . . .                   | 52        |
| 4.3      | PINN Solution . . . . .                               | 53        |
| 4.3.1    | Variational equation . . . . .                        | 53        |
| 4.3.2    | PINN formulation . . . . .                            | 54        |
| 4.4      | Multi-asset PINN . . . . .                            | 55        |
| 4.4.1    | One-dimensional reduction . . . . .                   | 56        |
| 4.4.2    | Variational equation . . . . .                        | 58        |
| 4.4.3    | PINN formulation . . . . .                            | 59        |
| 4.5      | Numerical Results . . . . .                           | 60        |
| 4.5.1    | Interpolating tree method . . . . .                   | 60        |
| 4.5.2    | One-dimensional PINN solution . . . . .               | 61        |
| 4.5.3    | Multi-asset PINN solution . . . . .                   | 62        |
| 4.5.4    | Free boundary . . . . .                               | 63        |
| <b>5</b> | <b>Conclusion</b>                                     | <b>66</b> |

# List of Figures

|      |                                                                   |    |
|------|-------------------------------------------------------------------|----|
| 1.1  | Option value illustration . . . . .                               | 9  |
| 1.2  | Volatility smile illustration . . . . .                           | 10 |
| 2.1  | Exercise boundary illustration . . . . .                          | 15 |
| 2.2  | Binomial tree illustration . . . . .                              | 19 |
| 3.1  | Asset product price histogram . . . . .                           | 31 |
| 3.2  | Tree convergence at $t = 0, S = K$ . . . . .                      | 36 |
| 3.3  | Tree vs closed-form: varying $S$ . . . . .                        | 36 |
| 3.4  | Tree vs closed-form: varying $t$ . . . . .                        | 37 |
| 3.5  | 1D PINN loss convergence . . . . .                                | 37 |
| 3.6  | PINN vs tree heatmap . . . . .                                    | 38 |
| 3.7  | PINN vs tree: varying $S$ and term structure . . . . .            | 38 |
| 3.8  | Training epochs for $f(0, K)$ . . . . .                           | 39 |
| 3.9  | PINN vs tree: $S^1$ - $S^2$ heatmap at $t = 0$ . . . . .          | 39 |
| 3.10 | PINN vs tree: at the money and term structure . . . . .           | 40 |
| 3.11 | Sobolev PINN vs tree heatmap . . . . .                            | 40 |
| 3.12 | Sobolev PINN vs tree: varying $S$ and term structure . . . . .    | 41 |
| 3.13 | Sobolev PINN vs tree: price heatmap . . . . .                     | 41 |
| 3.14 | Sobolev PINN vs tree: varying $S$ and term structure . . . . .    | 42 |
| 3.15 | PINN continuation probabilities (1D) . . . . .                    | 43 |
| 3.16 | Price estimate distribution (25 runs) . . . . .                   | 43 |
| 3.17 | 2D continuation probabilities: $S^1 \times S^2 = K$ . . . . .     | 44 |
| 3.18 | 2D continuation probabilities: $S^1 = S^2$ slice . . . . .        | 44 |
| 3.19 | Put on minimum of two assets heatmap . . . . .                    | 45 |
| 3.20 | Put on minimum of two assets continuation probabilities . . . . . | 46 |
| 4.1  | Heston Volatility Smile . . . . .                                 | 48 |
| 4.2  | Heston Tree interpolation illustration . . . . .                  | 51 |
| 4.3  | Heston tree call option prices . . . . .                          | 60 |
| 4.4  | Heston tree: American vs European put options . . . . .           | 61 |

|      |                                                                     |    |
|------|---------------------------------------------------------------------|----|
| 4.5  | Heston 1D PINN vs tree comparison heatmap . . . . .                 | 61 |
| 4.6  | Heston 1D PINN vs tree varying $S$ , $V$ and $t$ . . . . .          | 62 |
| 4.7  | Heston nD PINN vs tree comparison heatmap . . . . .                 | 62 |
| 4.8  | Heston PINN continuation probabilities (1D) . . . . .               | 63 |
| 4.9  | Heston PINN continuation probabilities (nD) $S^1$ - $S^2$ . . . . . | 64 |
| 4.10 | Heston PINN continuation probabilities (nD) $S$ - $t$ . . . . .     | 65 |



# Chapter 1

## Introduction


### 1.1 Motivation

In addition to bonds and stocks, the financial market offers derivatives, which are financial instruments whose value depend on the value of an underlying asset, which could be a stock, a bond, an index or even another derivative. These are useful in the real world for hedging and speculation, widely used by institutions and individuals alike, and are mathematically rich objects that have been the subject of intense research for decades.

Options are a specific type of derivative contract that gives the holder the right, but not the obligation, to buy or sell an underlying asset at a predetermined strike price  $K$ , specified at the time of option purchase. These contracts have a fixed expiry time  $T$ , at which the holder must decide whether to exercise the option or let it expire worthless. Different types of options exist: European options can only be exercised at the expiry time  $T$ , while American options can be exercised at any time before expiration. With the increased flexibility of American options, they are more widely used in practice, and are more interesting mathematical objects to study, since they are much more difficult to price. Holders of this contract must ask themselves constantly whether they should settle for the current value by exercising the option immediately, or wait for a potentially better opportunity.

In this project, we will concentrate our attention towards put options (which give the holder the right to sell the underlying asset), since for call options (which give the holder the right to buy the underlying asset), the American option is equivalent to the European option, since we do not consider dividends in this project. At expiry, the payoff of the put option is given by  $\max(K - S_T, 0)$ , where  $S_T$  is the price of the underlying asset at time  $T$ . This is because holders of the put option should sell when the price of the underlying asset is below the strike price. Similarly, for a call option, holders should buy when the price  $S_T > K$ , and hence the payoff is given by  $\max(S_T - K, 0)$ . If the underlying asset price is such that we should exercise the option, the option is in-the-money (expires with a positive payoff). Else, the option is out-of-the-money (expires worthless). Figure 1.1 illustrates these payoff functions at expiry with the grey line.

Figure 1.1 also shows the value of a put and a call option at the initial time, which can be decomposed into the intrinsic value (the payoff should it be exercised immediately) and the time value (the additional value from the possibility of future price movements). We can see clearly that the European option price bounds the American option price from below, since the holder of the European option cannot exercise early. Interestingly, we can see that the value of the European put is lower than the intrinsic value for in-the-money asset prices, which is a feature of the positive interest rate coupled with the fact the option can only be exercised at expiry. Of course, the intrinsic value bounds the American option from below, since the holder



**Figure 1.1** Option values (for put and call) at initial time


can always choose to exercise immediately.

These options are very important financial instruments, and tens of millions of contracts are traded every day just for the S&P 500 index alone. For that reason, it is important to be able to price these options accurately. This is a non-trivial task, as the price of an option depends on many factors, such as the underlying asset price which is generally modelled as a stochastic process, and the American option does not admit a closed-form solution.

## 1.2 Literature review

The pioneering work in option pricing was done by Black and Scholes [1], who derived a closed-form solution for the price of the European call option under the assumption that the underlying asset price follows a geometric Brownian motion. This model, known as the Black-Scholes model, is considered the foundational work in option pricing theory, and the authors were awarded the Nobel Prize in Economics in 1997 for their contribution. However, it relies on several assumptions such as constant risk-free rate and volatility, which is not always realistic in practice, with real markets showing implied volatility varying across different strike prices and maturities, a phenomenon known as the volatility smile (see Figure 1.2). As a result, people started developing more sophisticated models that took inspiration from the Black-Scholes model. Heston [2] modelled volatility to be a stochastic process in itself, which is more consistent with empirical observations of financial markets, and derived the European option prices for this proposed model. This allowed a better fit to the volatility smile and was more flexible in capturing market dynamics.

While the European option has an analytical solution for specific, well-researched models, the American option is much more difficult to price, since it allows the holder to exercise the option at any time before expiration, making it analogous to an optimal stopping problem. These optimisation problems do not admit a closed-form solution in general. As a result, various numerical methods have been developed to tackle this problem, such as the binomial tree method [3, 4] for the Black-Scholes and Heston models respectively. In these methods, the asset price is discretised into a tree structure and the option prices are computed via backwards induction from the terminal payoff at expiry. Monte-Carlo simulation methods were also explored [5], where heavy research has gone into optimising the idea of pricing an option by simulating an ensemble of paths and computing the average value. More recently, variational methods were



**Figure 1.2** Volatility smile for SPX options with expiry on 30 June 2026. Data from Yahoo Finance (retrieved on 1 June 2026).

explored, where the optimal stopping problem is formulated as a variational inequality, with deep neural networks finding approximations to the solution [6, 7], borrowing from the field of machine learning and physics. This method is further adapted to incorporate fractional Sobolev norms to guarantee convergence of the solution to the true option price instead of just pointwise convergence, which has been particularly developed for American option pricing [8].

### 1.3 Overview of project

This project aims to numerically solve the American option pricing problem under the Black-Scholes and Heston models on increasingly complex payoffs by iteratively building benchmarks.

For each individual model, we begin by deriving the price of a European call option, which serves as a benchmark for the binomial tree method. These formulae are well-known and can be found in the original papers [1, 2] but we provide a condensed derivation for completeness, with clearer clarification on the results employed in the steps in the derivation.

The binomial trees can easily be adapted to price American options, which provide a benchmark for the neural network solution. The binomial tree method is a well-established numerical technique for pricing options, but we take an interpolation approach, as suggested in [4], for the Heston model where there are two concurrent processes (the asset and the volatility) to build a tree for. The original method yields good results, but we further adapt the method to allow efficient pricing across the whole pricing domain.

The neural network method is implemented with original loss functions inspired by the PINN framework based on the boundary behaviour of the option price, with weights tuned by adaptive weighting schemes. We additionally implement an alternative loss function based on the fractional Sobolev norm, an original research field explored in [8] by one of my supervisor's PhD students, which we adapt to price American options under the whole domain space instead of just at-the-money as in the original paper.

We use both the neural network and the binomial tree method as benchmarks for the multi-asset case, where we price put options on the product of assets, which has yet to be explored in the literature in depth. This is motivated by the fact that there exists an equivalent single-asset reduction for this payoff, which allows us to validate the results against the benchmarks already

established in the one-asset case. With this foundation established, we are then free to tackle alternative payoffs where such a reduction does not exist, such as the maximum or minimum of several assets, and still have confidence in the results since the methods have been validated on the simpler cases.

We will conclude our numerical results by applying our pricing scheme to an estimation of the free boundary. The free boundary is the critical asset price at which the holder of the option is indifferent between exercising immediately or waiting, which we compute numerically through a binary search method for deterministic methods such as the binomial tree. For non-deterministic methods, we propose a novel method to frame the free boundary problem as a soft-classification problem, where we make distributional assumptions on our continuation values and use the empirical variance of the continuation values as a measure of uncertainty that tunes the soft classification. The continuation values are estimated through Monte-Carlo simulation, and the soft classification is done through an empirically tuned injective function.

The project aims to provide a comprehensive numerical framework for pricing American options under two different models, with detailed numerical results highlighting the strengths and shortcomings of each method, and insights into the option price behaviour under different payoffs, complexity and market conditions. All of these are backed by a rigorous mathematical framework that borrows from the fields of stochastic calculus, numerical analysis and machine learning, allowing a translation of theoretical insights into practical numerical methods that can be applied in real-world scenarios.

All of the methods explored in this project are implemented in Python, with the code available on [GitHub](#)<sup>1</sup>.

---

<sup>1</sup>This repository can be accessed with the url [github.com/jameswu5/M4R](https://github.com/jameswu5/M4R).

## Chapter 2

# Mathematical Background

### 2.1 Preliminaries

We will first assemble the necessary mathematical background to understand the option pricing problem and build some foundations into the methods we will be using to solve it. The definitions and results are mostly standard, and can be found in textbooks such as [9, 10] but we will cover them here for self-containment and to set up notation for the rest of the project. As a result, we will only cover the main concepts and not delve too deeply into the details, since our focus is more on the numerical methods and their applications.

#### 2.1.1 Probabilistic Setup

The option pricing problem is inherently probabilistic, as we are uncertain about the future evolution of the underlying asset price. Therefore, we need a rigorous mathematical framework to model this uncertainty. The formalisation of the underlying randomness in our model for the market will be done through *probability spaces*, and the arrival of information in the market through time through *filtrations*.

**Definition 2.1.1** (Probability space). A *probability space* is a triple  $(\Omega, \mathcal{F}, \mathbb{P})$  where  $\Omega$  is the sample space,  $\mathcal{F}$  is a  $\sigma$ -algebra on  $\Omega$ , and  $\mathbb{P} : \mathcal{F} \rightarrow [0, 1]$  is a probability measure such that  $\mathbb{P}(\Omega) = 1$ .

**Definition 2.1.2** (Filtration, filtered probability space). A *filtration* on the probability space  $(\Omega, \mathcal{F}, \mathbb{P})$  is a non-decreasing family  $(\mathcal{F}_t)_{t \geq 0}$  of sub- $\sigma$ -algebras of  $\mathcal{F}$ , that is to say,  $\mathcal{F}_s \subseteq \mathcal{F}_t \subseteq \mathcal{F}$  for all  $s \leq t$ . The quadruple  $(\Omega, \mathcal{F}, \{\mathcal{F}_t\}_{t \geq 0}, \mathbb{P})$  is called a *filtered probability space*.

We will also define a Martingale, which formalises the idea that future expectations of a process are equal to the current value (there is no drift in the process), which is a key concept in the risk-neutral pricing framework.

**Definition 2.1.3** (Martingale). A process  $M = (M_t)_{t \geq 0}$  is a *martingale* with respect to the filtration  $\{\mathcal{F}_t\}$  and the probability measure  $\mathbb{P}$  if  $\mathbb{E}[M_t | \mathcal{F}_s] = M_s$  for all  $0 \leq s \leq t$ .

The Markov assumption is also commonly made in financial modelling, which states that the future evolution of the process only depends on the current state and not on the past history. It is useful because it greatly simplifies models, and hence we will employ it in our project as well.

**Definition 2.1.4** (Markov process). A process  $X = (X_t)_{t \geq 0}$  is a *Markov process* with respect to the filtration  $\{\mathcal{F}_t\}$  if

$$\mathbb{E}[f(X_t) | \mathcal{F}_s] = \mathbb{E}[f(X_t) | X_s] \quad (2.1)$$

for all  $0 \leq s \leq t$  and all bounded measurable functions  $f$ .

### 2.1.2 Stochastic processes

The market in practice is not deterministic, so we will mathematically construct our model with stochastic processes, indexed by time in our application.

**Definition 2.1.5** (Stochastic process). A *stochastic process* is a family of random variables  $\{X_t : t \geq 0\}$  defined on a probability space  $(\Omega, \mathcal{F}, \mathbb{P})$ .

More specifically, we will use Brownian motion to model the underlying asset price dynamics, which is a continuous-time stochastic process with continuous paths and stationary independent increments.

**Definition 2.1.6** (Brownian motion). A real-valued process  $W = (W_t)_{t \geq 0}$  on the probability space  $(\Omega, \mathcal{F}, \mathbb{P})$  is a *Brownian motion* if it satisfies the following properties:

- $W_0 = 0$  almost surely
- The increment  $W_t - W_s$  is normally distributed with mean 0 and variance  $t - s$  for all  $0 \leq s < t$
- The increments  $W_{t_1} - W_{t_0}, \dots, W_{t_n} - W_{t_{n-1}}$  are independent for all  $0 \leq t_0 < t_1 < \dots < t_n$
- The sample paths  $t \mapsto W_t(\omega)$  are continuous

The existence of this process is established by Wiener [11] and is very powerful in modelling the randomness in the financial market.

We will now define the Itô process [12], which is a generalisation of Brownian motion, allowing for a drift term and a diffusion term, giving us more flexibility to model the asset price dynamics.

**Definition 2.1.7** (Itô process). Suppose  $W = (W^1, \dots, W^d)$  is a  $d$ -dimensional Brownian motion on a filtered probability space  $(\Omega, \mathcal{F}, \{\mathcal{F}_t\}, \mathbb{P})$ . An *Itô process* is a process  $X = (X^1, \dots, X^d)$  valued in  $\mathbb{R}^d$  such that

$$X_t = X_0 + \int_0^t \mu_s ds + \int_0^t \sigma_s dW_s \quad (2.2)$$

where  $\mu \in \mathbb{R}^n$  and  $\sigma \in \mathbb{R}^{n \times d}$  are adapted processes satisfying some integrability conditions. We can also write this in differential form as

$$dX_t = \mu_t dt + \sigma_t dW_t \quad (2.3)$$

Finally, we present Itô's lemma [12], which allows us to compute the differential of a function of an Itô process, which we will use extensively in deriving the PDEs in the option pricing problem.

**Theorem 2.1.1** (Itô's lemma). Let  $X = (X^1, \dots, X^d)$  be an Itô process and  $f : \mathbb{T} \times \mathbb{R}^d \rightarrow \mathbb{R} \in \mathcal{C}^{1,2}$  be a function that is once continuously differentiable in time and twice continuously differentiable in space. Then the process  $f(t, X_t)$  is also an Itô process and satisfies

$$\begin{aligned} f(t, X_t) = f(0, X_0) &+ \int_0^t \frac{\partial f}{\partial t}(s, X_s) ds + \sum_{i=1}^d \int_0^t \frac{\partial f}{\partial x_i}(s, X_s) dX_s^i \\ &+ \frac{1}{2} \sum_{i=1}^d \sum_{j=1}^d \int_0^t \frac{\partial^2 f}{\partial x_i \partial x_j}(s, X_s) d\langle X^i, X^j \rangle_s \end{aligned} \quad (2.4)$$

### 2.1.3 Risk-neutral pricing

We will now formalise the financial market and establish the principle of arbitrage-free valuation that will underpin our option pricing problem.

In our model, the market involves a fixed horizon  $T > 0$  (our expiry) and a filtered probability space  $(\Omega, \mathcal{F}, \{\mathcal{F}_t\}_{t \in [0, T]}, \mathbb{P})$  where  $\{\mathcal{F}_t\}$  is the filtration generated by the underlying Brownian motion. The market consists of a risk-free asset  $B_t$  that is deterministic and follows the evolution

$$B_t = \exp \left( \int_0^t r(s) ds \right) \quad (2.5)$$

and one or more risky assets  $S_t$  that follows an Itô process.

It is important that the market is arbitrage-free, which means that there does not exist any trading strategy that can generate a positive profit with zero initial investment and zero risk. This trading strategy  $\phi$  must be self-financing, which means that any change in the value of the portfolio is only due to changes in the assets held, without any external cash flow.

**Definition 2.1.8** (Self-financing trading strategy). A *self-financing trading strategy*  $\phi$  is a predictable process such that the value of the portfolio  $V_t = \phi_t^B B_t + \phi_t^S S_t$  satisfies

$$dV_t = \phi_t^B dB_t + \phi_t^S dS_t \quad (2.6)$$

**Definition 2.1.9** (Arbitrage). An *arbitrage* is some trading strategy  $\phi$  such that  $V_0 = 0$ ,  $V_T \geq 0$  almost surely and  $\mathbb{P}(V_T > 0) > 0$ .

By the fundamental theorem of asset pricing, the market is arbitrage-free if and only if there exists a measure  $\mathbb{Q}$  such that  $\mathbb{Q}$  is equivalent to  $\mathbb{P}$  (they have the same null sets) and the discounted price process  $\tilde{S}_t = \frac{S_t}{B_t}$  is a martingale under  $\mathbb{Q}$ . We will assume that our market is complete (any derivative can be replicated), and hence such a measure  $\mathbb{Q}$  will be unique.

Given that the market is arbitrage-free, we can price any derivative with payoff  $g(\cdot)$  at time  $t$  by taking the discounted expectation of the payoff under the risk-neutral measure  $\mathbb{Q}$ , giving us the risk-neutral pricing formula for European derivative with maturity  $T$ :

$$v(t, S) = B_t \mathbb{E}^{\mathbb{Q}} \left[ \frac{g(S_T)}{B_T} \mid \mathcal{F}_t \right] \quad (2.7)$$

Throughout this project specifically, we will deal with constant interest rates, so the bond price process reduces to  $B_t = e^{rt}$  where  $r$  is the constant risk-free rate. Additionally, we will also assume the Markov property (2.1.4). The arbitrage-free price of the European derivative will therefore be given by

$$v(t, S) = \mathbb{E}^{\mathbb{Q}} \left[ e^{-r(T-t)} g(S_T) \mid S_t = S \right] \quad (2.8)$$

### 2.1.4 Optimal stopping problem and free boundary

The American option gives the holder the right to exercise the option at any time before maturity, which thus gives us the interpretation of the option price as the solution to an optimal stopping problem, as we need to additionally identify when we should exercise the option.

**Definition 2.1.10** (Stopping time). A random variable  $\tau : \Omega \rightarrow [0, \infty]$  is a *stopping time* if  $\{\tau \leq t\} \in \mathcal{F}_t$  for all  $t \geq 0$ .

The price of the American option can thus be expressed as the  $\mathbb{Q}$ -expectation of the discounted payoff at the optimal stopping time  $\tau^*$  that maximises the expected payoff, giving the optimal stopping problem

$$v(t, S) = \sup_{\tau \in [t, T]} \mathbb{E}^{\mathbb{Q}} \left[ e^{-r(\tau-t)} g(S_{\tau}) \mid S_t = S \right] \quad (2.9)$$

which is simply (2.8) with the additional optimisation over the stopping time  $\tau$ .

This optimal  $\tau^*$  is the first time that the option holder chooses to exercise the option, which motivates the notion of a free boundary, which is the boundary separating the exercise region and the continuation region.

Let's work with the case  $t = 0$  for simplicity, but the same logic applies for any  $t \in [0, T)$ . For the American put, the optimal stopping problem (2.9) can be rewritten as

$$v(0, S) = \sup_{\tau \in [0, T]} \mathbb{E} \left[ e^{-r\tau} (K - S_{\tau})^+ \mid S_0 = S \right] \quad (2.10)$$


which is achieved by an optimal stopping time  $\tau^*$  taking the form

$$\tau^* = \inf \{ \tau \geq 0 : S_{\tau} \leq b^*(\tau) \} \quad (2.11)$$

for some optimal exercise boundary  $b^*(\tau)$  [5]. This is illustrated in Figure 2.1, where we simulate one realisation of the asset price assuming a geometric Brownian motion and plot it against the exercise boundary. The illustration is based on the binomial tree solution on the American put option, which has the boundary

$$b^*(\tau) = \sup \{ S : c(\tau, S) = g(S) \} \quad (2.12)$$

which is found using binary search here, and  $c(t, S)$  is the continuation value of the option with payoff  $g(S)$ .



**Figure 2.1** Exercise boundary for the American put (adapted from [5])



## Soft classification of the free boundary

This formulation describes a hard classification problem, where we classify points in the time-stock price space binarily to be either in the continuation region or the exercise region. However, our numerical solutions to the pricing problem are not exact and there is some level of error in the solution that comes from the Monte Carlo simulations upon the space, the numerical differentiation used in training and the early stopping in the training process. Therefore, it would be more intuitive to perform a soft classification instead, where we assign probabilities to points in the time-stock price space to be in the continuation region or the exercise region, which would give us a more intuitive visualisation of the free boundary and how it evolves over time. We will do this by extending hard classification: that is to say, the continuation region  $A$  is

$$A = \{(t, S) : c(t, S) > g(S)\} \cup \{(t, S) : g(S) = 0\} \quad (2.13)$$

where intuitively, the first set is those where the continuation value is greater than the early exercise value (and hence it is better to continue), and the second set is those where the payoff is zero, which are at-the-money and out-of-the-money options that are worthless under immediate exercise.

## Sigmoid soft classification

We will consider two methods to extend the definition of the continuation region to a soft classification. The first one is to use a sigmoid function  $s : \mathbb{R} \rightarrow (0, 1)$  defined as

$$s(x) = \frac{1}{1 + e^{-x}} \quad (2.14)$$

to smoothly transition the continuation probability from 0 to 1 as the value function transitions from being below the payoff to being above the payoff. As the sigmoid function is the CDF of the logistic distribution, we can interpret this as modelling the noise in the value function as a logistic distribution. The effect of the noise can be controlled by introducing the softness parameter  $\tau > 0$ , and our continuation probability function can be defined on the relative difference from the payoff as

$$P_{\text{cont}}(d) = s\left(\frac{d}{\tau}\right), \quad d := \frac{c(t, S) - g(S)}{g(S)} \quad (2.15)$$

To tune the softness parameter, we can define a threshold  $I \in (0.5, 1)$  such that we want the continuation probability to be  $I$  at some relative difference  $\epsilon > 0$ , i.e.  $P_{\text{cont}}(\epsilon) = I$ , which we could solve with some simple algebraic manipulation to get

$$\tau = \frac{\epsilon}{\ln\left(\frac{I}{1-I}\right)} \quad (2.16)$$

## Gaussian noise soft classification

The second method is to assume Gaussianity of the noise in our estimate  $\hat{c}$  of the continuation value, and that our approximated price is unbiased. In other words, we have

$$\hat{c}(t, S) = c(t, S) + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \tilde{\sigma}^2) \quad (2.17)$$

where  $c$  is the true (latent) continuation value and  $\tilde{\sigma}$  is the standard deviation of the solution noise, which is also unknown. We can then perform solve the problem  $n$  times to get the

different approximated continuation values  $\hat{c}_1, \dots, \hat{c}_n$ . This allows us to use the sample mean  $\bar{c}$  as our unbiased estimate of the continuation value function paired with the sample standard deviation  $\hat{\sigma}$  as our estimate. Then  $\bar{c}$  is distributed according to a  $t$ -distribution with  $n - 1$  degrees of freedom, and we can compute the continuation probability as

$$P_{\text{cont}}(t, S) = \mathbb{P}_{t_{n-1}} \left( \frac{\bar{c}(t, S) - g(S)}{\hat{\sigma}(t, S)/\sqrt{n}} > 0 \right) \quad (2.18)$$

which is simple as the CDF of the  $t$ -distribution is readily available.

The caveat with this method is that it requires multiple simulations of the solution, which becomes infeasible for more complex models and higher dimensions, but it gives us a more statistically grounded way to compute the continuation probability.

All that is left is to find a method to obtain the continuation value function, as it is generally latent and not directly observable, given that we predict the American option price that already takes the maximum of the continuation value and the payoff. However, we can use

$$c(t, S) = \mathbb{E} \left[ e^{-rh} f(t + h, S_{t+h}) \mid S_t = S \right] \quad (2.19)$$

where  $h > 0$  is a small time step and  $f$  is the American option price function. We can obtain an unbiased estimation of this expectation by Monte Carlo: we simulate the asset price  $S_{t+h}$  at time  $t + h$  given  $S_t = S$  many times and evaluate the option price function for each simulation, then taking the mean.

### 2.1.5 Monte Carlo integration

Monte Carlo is a powerful way to estimate integrals that may be intractable or not have a closed-form solution, which will be useful through our project as expectations can be expressed as integrals under the probability measure. The idea is to use the law of large numbers to estimate the integral by sampling from some distribution and evaluating the integrand at the sampled points.

Suppose we want to compute the integral of some function  $\phi$  over a domain  $\Omega$ :

$$I = \int_{\Omega} \phi(x) \, dx \quad (2.20)$$

If we are able to sample iid from some distribution with probability density function  $q$  that has positive density on  $\Omega$ , then we can rewrite the integral as

$$I = \int_{\Omega} \frac{\phi(x)}{q(x)} q(x) \, dx = \mathbb{E} \left[ \frac{\phi(X)}{q(X)} \right] \quad (2.21)$$

where the expectation is taken under the distribution with density  $q$ . This expectation can be estimated by sampling  $X_1, \dots, X_N$  iid from  $q$  and taking the sample mean:

$$\hat{I} = \frac{1}{N} \sum_{j=1}^N \frac{\phi(X_j)}{q(X_j)} \quad (2.22)$$

This is known as importance sampling [13], and our choice of  $q$  can affect the variance of the estimator  $\hat{I}$ . In practice, we can choose  $q$  to be uniform over  $\Omega$ , giving us the simple estimator

$$\hat{I} = \frac{1}{N} \sum_{j=1}^N \phi(X_j) \quad (2.23)$$

## 2.2 State Space Discretisation

### 2.2.1 Background

Recall that the option value function  $v(t, S)$  is the solution to the optimal stopping problem

$$v(t, S) = \sup_{\tau \in [t, T]} \mathbb{E} \left[ e^{-r(\tau-t)} g(S_\tau) \mid S_t = S \right] \quad (2.24)$$

There does not exist a closed-form solution for this problem in general, and thus we have to resort to numerical methods. We first consider methods that discretise the continuous time into a finite grid  $\{t_i\}_{i=1}^N$  where we have

$$0 = t_0 < t_1 < \dots < t_N = T \quad (2.25)$$

which we can use to index the underlying state variable  $S_{t_i}$  at time  $t_i$ . Note that the state variable  $S_{t_i}$  could depend on multiple variables, such as in the Heston model, where we have a stochastic variance term encoded in the asset price. The goal is to compute the arbitrage-free price  $v(0, S_0)$  at time  $t = 0$  given the initial state  $S_0$ . As we know the solution must satisfy the terminal condition  $v(T, S) = g(S)$  for all  $S$ , we can use backward induction to compute the option value at each time step. This gives rise to the class of numerical methods known as tree methods.

### 2.2.2 Tree Methods

In a tree model, we discretise the time-state space into a finite grid, where the state variable  $S_{t_i}$  can take a finite number of values at each time step given the previous one. We will particularly explore binomial trees in this project, where each state can evolve to one of two possible states at the next time step. Formally, suppose at time  $t_n$  we have the state  $S_{t_n} \in \{S_n^1, \dots, S_n^{M_n}\}$  where  $M_n$  is the number of possible states at time  $t_n$ . Then at time  $t_{n+1}$ , each state  $S_n^i$  can evolve to the state  $S_{n+1}^j$  with transition probability

$$p_{ij}^{(n)} = \mathbb{P}(S_{n+1} = S_{n+1}^j \mid S_n = S_n^i) \quad (2.26)$$

This is a discrete-time Markov chain, and to identify the transition probabilities, we can use the risk-neutral condition which states that the discounted price process must be a martingale under the risk-neutral measure. This gives us the condition

$$\sum_{j=1}^{M_{n+1}} p_{ij}^{(n)} S_{n+1}^j = e^{r\Delta t} S_n^i \quad (2.27)$$

where  $\Delta t = t_{n+1} - t_n$  and  $r$  is the risk-free rate.


With these transition probabilities, we can compute the option value at time  $t_n$  given the option value at time  $t_{n+1}$  by using backwards induction, which gives us the recursive relation

$$v(t_n, S_n^i) = \max \left( g(S_n^i), e^{-r\Delta t} \sum_{j=1}^{M_{n+1}} p_{ij}^{(n)} v(t_{n+1}, S_{n+1}^j) \right) \quad (2.28)$$

where we take the maximum of the payoff and the discounted expected future price to account for early exercise for the American option. This maximum operator can easily be lifted for all

time steps for European options, or only imposed on certain time steps for Bermudan options, leaving the tree method a very flexible method. This backwards induction is applied iteratively starting from the terminal condition at time  $t_N = T$  to compute the option value at time  $t_0 = 0$ .

An issue with tree methods is that the number of states  $M_n$  can grow exponentially with the number of time steps  $N$ , leading to a very computationally expensive algorithm. To mitigate this, we use recombining trees where the states at each time step are designed to recombine with each other, so that the number of states  $M_n$  grows linearly. Figure 2.2 illustrates the difference between non-recombinant and recombinant trees. Another issue is that, while flexible, the tree method can be difficult to implement for more complex models where the state depends on multiple variables or if we wish to extend it to options operating on multiple assets, but its reliability and interpretability make it a good benchmark for simpler models.



**Figure 2.2** Non-recombinant trees vs recombinant trees. The number of nodes in non-recombinant trees scales exponentially with the number of time steps.

## 2.3 Physics-Informed Neural Networks

### 2.3.1 Motivation

Our option pricing problem can alternatively be formulated as a variational inequality which takes the form of a partial differential equation that is independent from the payoff  $g(\cdot)$ . These PDEs can be of high dimension, and are thus difficult to solve using traditional numerical methods due to the curse of dimensionality. Traditional methods such as finite difference methods require discretisation of the domain and the number of grid points (and hence algorithm complexity) scales exponentially  $\mathcal{O}(N^d)$  with the dimension  $d$  of the problem.

We can attempt a neural network solution to the PDEs, as neural networks are known to be universal approximators. That is, suppose we have the class of functions  $\mathcal{F}$  over a compact set  $S$  that can be represented by a neural network with a given architecture, with a continuous target function  $f^*$  that we would like to approximate. Then for any target accuracy  $\epsilon > 0$ , there exists a function  $f \in \mathcal{F}$  such that

$$\sup_{x \in S} |f(x) - f^*(x)| \leq \epsilon \quad (2.29)$$

If we have a squashing activation function  $\Psi(\cdot)$  that is non-decreasing and satisfies the limits

$$\Psi(-\infty) = 0, \quad \Psi(\infty) = 1 \quad (2.30)$$

then the class of functions  $\mathcal{F}$  that can be represented by a neural network with one hidden layer and activation function  $\Psi$  is dense in the space of continuous functions on  $S$  and thus is a universal approximator [14].

Since the option price function is continuous in the time-asset price domain, we can use a neural network to find an approximation. This is favourable because the time complexity scales with the number of parameters, which is not exponential. To do this, we will use the framework of physics-informed neural networks (PINNs) to solve the PDEs, which have been shown to be effective for solving high-dimensional PDEs [6], by avoiding mesh construction and leveraging the interpolation capability of neural networks [15]. The loss function of the neural network is designed to penalise deviations from the PDE as well as deviations from the boundary conditions. By training the neural network to minimise this loss function, we can obtain an approximation to the solution of the PDE. This method has been applied to option pricing problems in the literature, such as in [16] for the Black-Scholes model, but we will look to apply it to more complex models such as the Heston model as well as to options on multiple assets.

### 2.3.2 Methodology

We first provide a general setup for PINNs, before modifying it to tailor our specific problem of option pricing. Formally, suppose we have an unknown latent solution  $u(t, x)$  which is governed by a PDE of the form

$$\frac{\partial u}{\partial t} + \mathcal{L}[u] = 0, \quad t \in [0, T], x \in \Omega \quad (2.31)$$

where  $\mathcal{L}$  is a (not necessarily linear) differential operator. Suppose additionally that this PDE is subject to the boundary conditions

$$u(0, x) = g(x), \quad x \in \Omega \quad (2.32)$$

$$\mathcal{B}[u] = 0, \quad t \in [0, T], x \in \partial\Omega \quad (2.33)$$

where  $\mathcal{B}$  is a boundary operator. Then the unknown solution  $u$  can be approximated by a deep neural network  $f_\theta(t, x)$  where  $\theta$  are the parameters of the neural network, trained by minimising a loss function composed of deviations from the PDE residual, boundary conditions and initial conditions [17]:

$$L(\theta) = \lambda_r L_r(\theta) + \lambda_{bc} L_{bc}(\theta) + \lambda_{ic} L_{ic}(\theta) \quad (2.34)$$

where the loss components are mathematically defined as

$$L_r(\theta) = \frac{1}{N_r} \sum_{i=1}^{N_r} |\mathcal{R}_\theta(t_r^i, x_r^i)|^2 \quad (2.35)$$

$$L_{bc}(\theta) = \frac{1}{N_{bc}} \sum_{i=1}^{N_{bc}} |\mathcal{B}[f_\theta](t_{bc}^i, x_{bc}^i)|^2 \quad (2.36)$$

$$L_{ic}(\theta) = \frac{1}{N_{ic}} \sum_{i=1}^{N_{ic}} |f_\theta(0, x_{ic}^i) - g(x_{ic}^i)|^2 \quad (2.37)$$

where the PDE residual  $\mathcal{R}_\theta$  is defined as

$$\mathcal{R}_\theta(t, x) = \frac{\partial f_\theta}{\partial t}(t, x) + \mathcal{L}[f_\theta](t, x) \quad (2.38)$$

We additionally use the sampled points  $\{x_{ic}^i\}_{i=1}^{N_{ic}}$  from the initial condition domain,  $\{(t_{bc}^i, x_{bc}^i)\}_{i=1}^{N_{bc}}$  from the boundary condition domain and  $\{(t_r^i, x_r^i)\}_{i=1}^{N_r}$  from the PDE domain, sampled uniformly in their respective domains at each iteration of the gradient descent algorithm.

There are different methods to choose the loss weights  $\lambda_r, \lambda_{bc}, \lambda_{ic}$ , such as using fixed weights or using adaptive weights that change during training, which we will discuss in Section 2.3.3.

For our specific problem of option pricing, we can modify the above setup to fit our variational inequality formulation. The PDE to solve is dependent on the model for the asset dynamics, and the boundary conditions are given by the payoff function of the option. In general, we will have to solve a problem of the form

$$\min \left\{ -\frac{\partial u}{\partial t} + \mathcal{L}[u], u - g \right\} = 0, \quad t \in [0, T], x \in \Omega \quad (2.39)$$

$$u(T, x) = g(x), \quad x \in \Omega \quad (2.40)$$

$$\mathcal{B}_k[u] = 0 \quad t \in [0, T], x \in \partial\Omega, k = 1, \dots, K \quad (2.41)$$

where  $k$  indexes the  $K$  different boundary conditions. This is very similar to the general setup for PINNs, except that we have a variational inequality instead of a PDE, and we have a terminal condition instead of an initial condition. However, we can still use the framework by noting we can make the substitution  $t \mapsto T - t$  to un-negate the negative partial derivative in the PDE term and to convert the terminal condition to an initial condition. The variational inequality can be handled by modifying the PDE residual loss (2.38) to

$$\mathcal{R}_\theta(t, x) = \min \left\{ \frac{\partial f_\theta}{\partial t}(t, x) + \mathcal{L}[f_\theta](t, x), f_\theta(t, x) - g(x) \right\} \quad (2.42)$$

With everything in place, we can train the neural network with gradient descent:

$$\theta_{n+1} = \theta_n - \eta_n \nabla_\theta L(\theta_n) \quad (2.43)$$

where  $\eta_n$  is the learning rate at iteration  $n$ , either kept constant or decayed over time with a schedule.

### 2.3.3 Adaptive Loss Weighting

While keeping the loss weights  $\lambda_r, \lambda_{bc}, \lambda_{ic}$  fixed is simple, it can be difficult to tune correctly and can lead to suboptimal performance. Ideally, we would like to balance the different loss components so that they are of the same order of magnitude, which can be achieved by using adaptive loss weighting methods, as proposed in [17]. The idea is to evaluate the contribution of each loss component on the total loss and adjust the weights accordingly to balance their contributions. In practice, we will use a moving average to prevent the weights from changing too rapidly, leading to instability in training. Mathematically, we compute the adaptive weights

$$\begin{aligned} \hat{\lambda}_r &= \frac{\|\nabla_\theta L_r(\theta)\| + \|\nabla_\theta L_{bc}(\theta)\| + \|\nabla_\theta L_{ic}(\theta)\|}{\|\nabla_\theta L_r(\theta)\|} \\ \hat{\lambda}_{bc} &= \frac{\|\nabla_\theta L_r(\theta)\| + \|\nabla_\theta L_{bc}(\theta)\| + \|\nabla_\theta L_{ic}(\theta)\|}{\|\nabla_\theta L_{bc}(\theta)\|} \\ \hat{\lambda}_{ic} &= \frac{\|\nabla_\theta L_r(\theta)\| + \|\nabla_\theta L_{bc}(\theta)\| + \|\nabla_\theta L_{ic}(\theta)\|}{\|\nabla_\theta L_{ic}(\theta)\|} \end{aligned}$$

and update the global weights  $\lambda = (\lambda_r, \lambda_{bc}, \lambda_{ic})$  with a smoothness  $\alpha \in (0, 1)$  as a moving average

$$\lambda_{\text{new}} = \alpha \hat{\lambda}_{\text{old}} + (1 - \alpha) \hat{\lambda}_{\text{new}} \quad (2.44)$$

This is not performed at each iteration, but rather at a fixed frequency to prevent the weights from changing too rapidly. The literature [17] suggests to update the weights every 1000 iterations with  $\alpha = 0.9$  in general.

## 2.3.4 Numerical Optimisation

### Gradient Descent

Our neural network will learn through gradient descent, in particular Adam, which uses adaptive learning rates and momentum to speed up convergence. Note this does not conflict with the adaptive loss weighting described in Section 2.3.3, as the adaptive loss weighting is used to balance the different loss components, while Adam is used to adapt the learning rate for each parameter based on the historical gradients.

We will additionally impose a schedule to decay the learning rate over time, which starts at a conventional value of  $10^{-3}$  to allow for fast learning at the beginning, and decays by a factor for every certain number of iterations, allowing for finer convergence in later iterations. These are also tuned by observing the training curve, and can vary for different models and different payoffs.

### Early Stopping

Early stopping is a regularisation technique mainly used to prevent overfitting in neural network training. However, we are training on a new random training set each epoch, so we do not need to worry about overfitting, and early stopping is implemented as an optimisation procedure.

The idea of early stopping is to stop the training process once we do not observe a significant improvement in the loss function after a certain number of iterations, which we call the patience. This can be tuned by observing the training curve and finding a good balance between training time and performance. Once the patience is exceeded without observing a significant improvement, the training is stopped and the model reverted to the best performing model observed during training. The loss used to evaluate the improvement is a validation set that is held through the training process, a set of points in the time-state space that is not used for training but only for evaluation.

Throughout this project, we will set the improvement threshold to be  $10^{-6}$  and the patience to be 500 iterations, which we found to be a good balance for the models we trained.

### Fine-tuning

The Adam optimiser is a first-order method that uses the first moment of the gradients to adapt the learning rate, so we do not leverage the information given by the second moment of the gradients. It is particularly good for navigating the loss landscape at the beginning of training, but in the later stages, it can exhibit poor convergence behaviour. This is because its stochastic noise and momentum can prevent it from tight convergence, requiring a very small learning rate to prevent oscillation. [18] explains in detail the shortcomings of Adam, where they construct explicit counterexamples of simple convex functions where Adam fails to converge to the optimal solution.

For PINNs, this is a problem, as we require a very small residual loss to satisfy the PDE accurately. We will therefore use a second-order method to fine-tune the model: Adam would serve as the global exploration to get us close to a minimum, and then we can switch to a second-order method to converge tightly. This makes us turn to quasi-Newton methods such as the BFGS algorithm [19], which maintains a dense Hessian matrix which is updated at each iteration, and a step size is computed with a line search method, which is more optimal than setting a learning rate. However, this can be computationally expensive for large models, and

the literature suggests we use L-BFGS [20] after the Adam training. This is a limited-memory version of BFGS that only stores a few vectors to approximate the Hessian, which is much more efficient for large models, at the cost of a shorter memory of the curvature information [21].



## Chapter 3

# Black-Scholes Model

The Black-Scholes model [1] is the fundamental model for option pricing, which assumes that the underlying asset price follows a geometric Brownian motion. As with the market introduced in Section 2.1.3, the market consists of two underlying assets:

- The risk-free asset, described by a deterministic function

$$dB_t = rB_t dt \quad (3.1)$$

where for convenience, we set  $B_0 = 1$  and  $r$  is the constant risk-free interest rate with continuous compounding. This has the solution  $B_t = e^{rt}$  for  $t \geq 0$ .

- The risky asset  $S_t$ , which has the price process of a geometric Brownian motion (GBM), which satisfies the stochastic differential equation

$$dS_t = \mu S_t dt + \sigma S_t dW_t \quad 0 \leq t \leq T \quad (3.2)$$

for a given drift  $\mu$  and volatility  $\sigma$ . We can solve this SDE to yield the unique solution

$$S_t = S_0 \exp \left( \left( \mu - \frac{1}{2} \sigma^2 \right) t + \sigma W_t \right) \quad 0 \leq t \leq T \quad (3.3)$$

For convenience, we will refer to the risk-free asset as the bond, and the risky asset as simply the asset for the remainder of this chapter.

Under risk-neutral pricing, we can find an equivalent probability measure  $\mathbb{Q}$  under which the discounted asset price is a martingale, and hence all assets earn the risk-free rate in expectation after discounting. Under the risk-neutral measure, the SDE for the asset becomes

$$dS_t = rS_t dt + \sigma S_t dW_t^{\mathbb{Q}} \quad (3.4)$$

where the drift has been replaced by the risk-free rate  $r$ . The derivation is done via Girsanov's theorem, and details can be found in Chapter 5 of [22].

### 3.1 Closed-form solution for European options

Our first benchmark will be the closed-form solution for the European option, which would be used to evaluate the accuracy of the binomial tree later. The derivation of this formula is widely

available and hence we will only include the main ideas. Using the risk-neutral pricing formula, the value of the European call option with strike  $K$  and expiry  $T$  is given by

$$v(0, S_0) = e^{-rT} \mathbb{E}^{\mathbb{Q}}[(S_T - K)^+ | S_0] \quad (3.5)$$

We use the solution of geometric brownian motion (3.3) to note that

$$S_T = S_0 \exp \left( (r - \frac{1}{2}\sigma^2)T + \sigma W_T^{\mathbb{Q}} \right) \quad (3.6)$$

under the risk-neutral measure  $\mathbb{Q}$ , having replaced the drift  $\mu$  with the risk-free rate  $r$ . The Brownian motion  $W_T^{\mathbb{Q}}$  has zero mean and variance  $T$ , so it can be written as  $\sqrt{T}Z$ , where we have the standard normal variable  $Z \sim \mathcal{N}(0, 1)$ . Therefore,  $S_T$  becomes

$$S_T = S_0 \exp \left( (r - \frac{1}{2}\sigma^2)T + \sigma\sqrt{T}Z \right) \quad (3.7)$$

This can be plugged into the risk-neutral pricing formula (3.5): the exercise event  $\{S_T > K\}$  is equivalent to  $\{Z > -d_2\}$  where

$$d_2 = \frac{\log(S_0/K) + (r - \frac{1}{2}\sigma^2)T}{\sigma\sqrt{T}}, \quad d_1 = d_2 + \sigma\sqrt{T} \quad (3.8)$$

The expectation (3.5) can be split, where we denote  $\mathbf{1}_A$  the indicator variable for the event  $A$ :

$$v(0, S_0) = e^{-rT} \mathbb{E}^{\mathbb{Q}}[(S_T - K)^+] \quad (3.9)$$

$$= e^{-rT} \mathbb{E}^{\mathbb{Q}}[(S_T - K)\mathbf{1}_{\{S_T > K\}}] \quad (3.10)$$

$$= e^{-rT} \mathbb{E}^{\mathbb{Q}}[S_T \mathbf{1}_{\{S_T > K\}}] - K e^{-rT} \mathbb{Q}(S_T > K) \quad (3.11)$$

The second probability is immediate as we see  $\mathbb{Q}(S_T > K) = \mathbb{Q}(Z > d_2) = \Phi(d_2)$  by symmetry, where  $\Phi$  is the CDF of the standard normal. We can complete the square in the Gaussian integrand for the first expectation to get the final formula [1]

$$v(0, S_0) = S_0 \Phi(d_1) - K e^{-rT} \Phi(d_2) \quad (3.12)$$

## 3.2 Binomial Tree Method

We can use the binomial discrete-time model to price both European and American options with the tree method. This model uses the simplification that the asset price can only move to one of two possible new prices in a time step. To avoid explosion in the number of nodes, we will only consider recombining trees so that the algorithm performs with quadratic time complexity (look back at Figure 2.2 for an illustration of recombining trees). We will verify the approximated solution's validity against the closed-form formula for European options, and use its American variant as a benchmark for later methods.

To formalise the tree construction, let's fix a time horizon  $T > 0$ . For some depth  $n \geq 1$ , divide the interval  $[0, T]$  into  $n$  even steps each of length  $\Delta t = T/n$ . Our asset  $S_t$ , where  $t$  is now a discrete time index, will follow a multiplicative binomial evolution:

$$S_{i+1} = \begin{cases} S_i u & \text{with probability } p \\ S_i d & \text{with probability } 1 - p \end{cases} \quad S_0 > 0 \quad (3.13)$$

where  $u$  and  $d$  are the model parameters determining the two factors of the asset price change per time step.

As with (3.1), the financial market will also consist of a risk-free asset  $B_t$ , discretised so that it satisfies

$$B_{i+1} = B_i e^{r\Delta t}, \quad B_0 = 1 \quad (3.14)$$

where  $r$  is the continuously-compounded risk-free rate, typically considered interest. For an arbitrage free market, we should additionally ensure that

$$d < e^{r\Delta t} < u \quad (3.15)$$

for sufficiently small  $\Delta t$ , else we can make a riskless profit by shorting the asset and buying the bond, or vice versa.

To employ risk-neutral pricing, we additionally require the discounted asset price to have zero drift (admitting a martingale measure). More precisely, let the discounted price be, with the bond as numeraire:

$$\tilde{S}_i = \frac{S_i}{B_i} = e^{-ri\Delta t} S_i \quad (3.16)$$

We require some probability measure  $\mathbb{Q}$ , equivalent to the physical measure  $\mathbb{P}$  (have the same non-zero support) such that

$$\mathbb{E}^{\mathbb{Q}}[\tilde{S}_{i+1} \mid \mathcal{F}_i] = \tilde{S}_i \quad (3.17)$$

where  $\mathcal{F}_i$  is the filtration at time step  $i$ .

To calculate the risk-neutral probability, we can apply this to our model:

$$e^{-r(i+1)\Delta t} [pS_i u + (1-p)S_i d] = e^{-ri\Delta t} S_i \quad (3.18)$$

Collect the  $p$  terms and cancel out the discount factor to get

$$p(S_i u - S_i d) + S_i d = e^{r\Delta t} S_i \quad (3.19)$$

Divide both sides by  $S_i$  and rearrange to make  $p$  the subject to obtain

$$p = \frac{e^{r\Delta t} - d}{u - d} \quad (3.20)$$

This is what we will use in our binomial model.

To price the option at time 0, we will perform backward induction from expiry, since we know that the price is just the payoff  $g : \mathbb{R}^+ \rightarrow \mathbb{R}$  as the option expires. As the discounted price process is a martingale under the risk-free measure  $\mathbb{Q}$ , the arbitrage-free value of the derivative at time  $i\Delta t$  is necessarily

$$v(t = i\Delta t) = \underbrace{e^{-r(T-i\Delta t)}}_{\text{discount}} \mathbb{E}^{\mathbb{Q}}[g(S_T) \mid \mathcal{F}_i] \quad (3.21)$$

We will utilise the Cox-Ross-Rubinstein (CRR) Model [3] for our implementation, which chooses the parameters

$$u = e^{\sigma\sqrt{\Delta t}} \quad d = e^{-\sigma\sqrt{\Delta t}} \quad (3.22)$$

satisfying the arbitrage free condition, and normalising the local variance of returns to the volatility parameter  $\sigma^2$ . They also conveniently multiply to 1.

To implement this, we need to first compute the price tree. This would be, for the price  $S_{i,j}$  at time step  $i$  after  $j$  up-moves (so clearly  $i \geq j$ )

$$S_{i,j} = S_0 u^j d^{i-j} \quad (3.23)$$

As the payoff is known at maturity, we can simply set, with  $v_{i,j}$  the value of the option at time step  $i$  with  $j$  up-moves:

$$v_{n,j} = g(S_{n,j}) \quad 0 \leq j \leq n \quad (3.24)$$

and recursively set

$$v_{i,j} = e^{-r\Delta t} (p v_{i+1,j+1} + (1-p) v_{i+1,j}) \quad (3.25)$$

for  $0 \leq j \leq i$ . This would be it for European options, but since early exercise is allowed, we will compare this value with the payoff of the price at step  $i$  with  $j$  up-moves, which is just  $f(S_{i,j})$ . Thus the value is given by

$$v_{i,j} = \max \left( e^{-r\Delta t} (p v_{i+1,j+1} + (1-p) v_{i+1,j}), g(S_{i,j}) \right) \quad (3.26)$$

The root value  $v_{0,0}$  would be the arbitrage-free option price.

### 3.3 Physics-Informed Neural Networks

In order to use the framework of PINNs as detailed in Section 2.3, we need to first derive the PDE satisfied by the option price, as well as the boundary conditions. We will first work with the one-dimensional case, then discuss the extension to the multi-dimensional case and how we could build a benchmark to against which we evaluate our results.

#### 3.3.1 PINN formulation for one-dimensional case

In order to set up our PINN, we need to first derive the PDE satisfied by the option price. Recall the value function  $v : [0, T] \times \mathbb{R}^+ \rightarrow \mathbb{R}$  of an American put with payoff  $g : \mathbb{R}^+ \rightarrow \mathbb{R}$  at time  $t$  with asset price  $S_t = S$  is given by

$$v(t, S) = \sup_{\tau \in [t, T]} \mathbb{E} [e^{-r(\tau-t)} g(S_\tau) \mid S_t = S] \quad (3.27)$$

where  $\tau$  is any stopping time in  $[t, T]$  and for convenience we will simply denote  $\mathbb{E}$  is conditional expectation at time  $t$ , given  $S_t = S$ , in this section. For an American option, we always have two options:

- Immediately exercise the option and get payoff  $g(S)$ . Since this is not necessarily optimal, the payoff must bound the value option from below:

$$v(t, S) \geq g(S) \quad (3.28)$$

- Do nothing and let the option continue to time  $t + h$  with  $h > 0$  a very small number. By using the supremum definition of  $v(t, S)$  (3.27) and swapping the expectation and supremum [23] we arrive at

$$e^{rh} v(t, S) \geq \mathbb{E} [v(t + h, S_{t+h})] \quad (3.29)$$

and the inequality turns into an equality at the limit  $\delta \rightarrow 0$ .

We will additionally assume that the value function  $v \in \mathcal{C}^{1,2}$  (continuously differentiable with respect to  $t$  and twice continuously differentiable with respect to  $S$ ) so that we can apply Itô's formula. Let's model the asset price  $S$  as a geometric Brownian motion satisfying under the risk-neutral measure

$$dS_t = rS_t dt + \sigma S_t dW_t \quad (3.30)$$

Then by Itô's lemma and taking expectations of both sides we have

$$\mathbb{E}[v(t + \delta, S_{t+\delta})] = v(t, S) + \mathbb{E}\left[\int_t^{t+\delta} \left(\frac{\partial v}{\partial t}(u, S_u) + \frac{\partial v}{\partial S}(u, S_u)rS_u + \frac{1}{2}\frac{\partial^2 v}{\partial S^2}(u, S_u)\sigma^2 S_u^2\right) du\right] \quad (3.31)$$

Now let's substitute (3.29) into the LHS, which gives us

$$e^{rh}v(t, S) \geq v(t, S) + \mathbb{E}\left[\int_t^{t+h} \left(\frac{\partial v}{\partial t}(u, S_u) + \frac{\partial v}{\partial S}(u, S_u)rS_u + \frac{1}{2}\frac{\partial^2 v}{\partial S^2}(u, S_u)\sigma^2 S_u^2\right) du\right] \quad (3.32)$$

We can use the Taylor approximation  $e^{rh} = 1 + rh + o(h)$ , which permits us to cancel the  $v(t, S)$  on both sides. Dividing both sides by  $h$  in the same step, we get

$$rv(t, S) + \frac{o(h)}{h} \geq \mathbb{E}\left[\frac{1}{h} \int_t^{t+h} \left(\frac{\partial v}{\partial t}(u, S_u) + \frac{\partial v}{\partial S}(u, S_u)rS_u + \frac{1}{2}\frac{\partial^2 v}{\partial S^2}(u, S_u)\sigma^2 S_u^2\right) du\right] \quad (3.33)$$

Sending  $h \rightarrow 0$  makes the higher order term  $o(h)/h$  vanish, and the expectation of the integral to collapse to the integrand. Thus we obtain

$$rv(t, S) \geq \frac{\partial v}{\partial t}(t, S) + \frac{\partial v}{\partial S}(t, S)rS + \frac{1}{2}\frac{\partial^2 v}{\partial S^2}(t, S)\sigma^2 S^2 \quad (3.34)$$

or for simplicity, we define the linear parabolic operator  $\mathcal{L}$  [8]:

$$\mathcal{L}[v] = rv(t, S) - \frac{\partial v}{\partial t}(t, S)rS - \frac{1}{2}\frac{\partial^2 v}{\partial S^2}(t, S)\sigma^2 S^2 \quad (3.35)$$

Thus we can concisely express this inequality as

$$\min\left(-\frac{\partial v}{\partial t}(t, S) + \mathcal{L}[v], v(t, S) - g(S)\right) \geq 0 \quad (3.36)$$

However, note that at time  $t$ , we must either exercise the option or do nothing and continue. One of these two options must be optimal, implying that one of the arguments must be 0. Thus we actually have an equality:

$$\min\left(-\frac{\partial v}{\partial t}(t, S) + \mathcal{L}[v], v(t, S) - g(S)\right) = 0 \quad (3.37)$$

which is the variational equation [24] for the American option. Note this is exactly the formulation of the PINN problem (2.39).

We are now ready to derive the boundary conditions. Our option has a terminal payoff at time  $T$  given by  $g(S)$ , so we have the terminal condition

$$v(T, S) = g(S) \quad (3.38)$$

which satisfies the form of the terminal condition for PINNs (2.40). For the spatial boundary conditions, we have two boundaries to consider:  $S \rightarrow 0$  and  $S \rightarrow \infty$ . In practice, we will have

to truncate the domain to  $\Omega = (S_{\min}, S_{\max})$  for some sufficiently small  $S_{\min}$  and sufficiently large  $S_{\max}$ . At the lower boundary, we have the boundary condition

$$v(t, S_{\min}) = g(S_{\min}) \quad (3.39)$$

because it is always optimal to exercise the option when the asset is at its lowest price. At the upper boundary, we assume the option is too far out of the money to ever finish in the money, and hence the option price is approximately 0. Thus we have the boundary condition

$$v(t, S_{\max}) = 0 \quad (3.40)$$

These three conditions together would form the boundary setup (2.41) for our PINN.

### 3.3.2 Multi-asset case

We will now extend our PINN framework to the multi-dimensional case, where we have  $n$  underlying assets, and hence the neural network  $f$  will take  $n + 1$  inputs (time and the price of each asset). The variational inequality can be derived with a similar argument as in the one-dimensional case, but we will use the multi-dimensional case of Itô's lemma (Theorem 2.1.1). In order to be able to evaluate the accuracy of our results, we will consider a specific multi-asset payoff such that there exists an equivalent one-dimensional reduction. This allows us to use the binomial tree method (or indeed the one-dimensional PINN solution) as a benchmark, giving us confidence for the PINN solution on more complex payoffs.

In the multi-dimensional case, we have  $n$  underlying assets with price processes following a correlated geometric Brownian motion. Mathematically, let each asset  $S^i$  satisfy under the risk-free measure  $\mathbb{Q}$ , for  $i = 1, \dots, n$ ,

$$dS_t^i = rS_t^i dt + \sigma_i S_t^i dW_t^i \quad (3.41)$$

with quadratic covariation  $d\langle W^i, W^j \rangle_t = \rho_{ij} dt$  for all  $i, j$  pairs.

### 3.3.3 Reduction of multi-asset to one-asset case

We will work with a specific payoff function that allows us to reduce the multi-dimensional case to a one-dimensional case, which would, in turn, give us confidence that our PINN implementation is correct, and we can change the payoff function to more complex ones and be reasonably assured that the results are still correct.

As the geometric Brownian motion is multiplicative, for our benchmark, we will consider the put on the product of the assets. More precisely, we consider the put option with payoff

$$g(S^1, \dots, S^n) = \left( K - \prod_{i=1}^n S^i \right)^+ \quad (3.42)$$

To represent the dynamics of the product of the assets, we can apply Ito's formula to the log of the price process. The log representation of the Itô process (refer to Definition 2.1.7) is

$$\ln S_t^i = \ln S_0^i + \int_0^t \left( r - \frac{1}{2} \sigma_i^2 \right) ds + \int_0^t \sigma_i dW_s^i \quad (3.43)$$

Define the new asset  $X$  to be

$$X_t := \prod_{i=1}^n S_t^i \quad (3.44)$$

and hence we have

$$\ln X_t = \sum_{i=1}^n \ln S_t^i = \ln X_0 + \int_0^t \left( nr - \frac{1}{2} \sum_{i=1}^n \sigma_i^2 \right) ds + \sum_{i=1}^n \int_0^t \sigma_i dW_s^i \quad (3.45)$$

By taking the differential we have

$$d \ln X_t = \left( nr - \frac{1}{2} \sum_{i=1}^n \sigma_i^2 \right) dt + \sum_{i=1}^n \sigma_i dW_s^i \quad (3.46)$$

The quadratic variation of the martingale term is

$$d \left\langle \sum_{i=1}^n \sigma_i W^i, \sum_{j=1}^n \sigma_j W^j \right\rangle_t = \sum_{i=1}^n \sum_{j=1}^n \sigma_i \sigma_j \rho_{ij} dt =: \sigma_X^2 dt \quad (3.47)$$

i.e.  $\sigma_X^2 = \sum_{i=1}^n \sum_{j=1}^n \sigma_i \sigma_j \rho_{ij}$ . So we can express the second term of (3.46) in terms of a new Brownian motion  $\widetilde{W}$ :

$$\sum_{i=1}^n \sigma_i dW_s^i = \sigma_X d\widetilde{W}_t \quad (3.48)$$

Going back, we have

$$d \ln X_t = \left( nr - \frac{1}{2} \sum_{i=1}^n \sigma_i^2 \right) dt + \sigma_X d\widetilde{W}_t \quad (3.49)$$

Exponentiating using Ito's formula gets us

$$dX_t = X_t \left[ \left( nr - \frac{1}{2} \sum_{i=1}^n \sigma_i^2 \right) dt + \sigma_X d\widetilde{W}_t \right] + \frac{1}{2} X_t \sigma_X^2 dt \quad (3.50)$$

which we can rearrange to get

$$dX_t = \left( nr - \frac{1}{2} \sum_{i=1}^n \sigma_i^2 + \frac{1}{2} \sigma_X^2 \right) X_t dt + \sigma_X X_t d\widetilde{W}_t \quad (3.51)$$

which is a geometric Brownian motion with volatility  $\sigma_X$  and drift  $r_X$  where

$$r_X = nr - \frac{1}{2} \sum_{i=1}^n \sigma_i^2 + \frac{1}{2} \sigma_X^2 \quad (3.52)$$

$$= nr + \sum_{i < j} \sigma_i \sigma_j \rho_{ij} \quad (3.53)$$


by matching coefficients. We can verify the equivalence by simulating many price paths for the assets and comparing their product against the simulations of price paths for the equivalent asset. Using the initial values (chosen arbitrarily)

$$r = 0.05 \quad S_0 = (10, 5, 3)^\top \quad \sigma = (0.2, 0.3, 0.1)^\top \quad \Sigma = \begin{pmatrix} 1 & 0.6 & -0.2 \\ 0.6 & 1 & 0.4 \\ -0.2 & 0.4 & 1 \end{pmatrix} \quad (3.54)$$

where  $\Sigma$  is the correlation matrix, we can plot the histogram of final prices in Figure 3.1.

They look very similar, and we can further statistically test if they come from the same distribution. Using Kolmogorov-Smirnov's 2-sample test on the null hypothesis that they come from the same distribution, we obtain a p-value of 0.8491, which does not provide sufficient evidence to reject the null hypothesis, which gives us stronger evidence that our derivation is correct.

Therefore, we will be able to price the put option on the product of multiple assets and compare that to the price of the put option on this single constructed asset as a benchmark.



**Figure 3.1** Histogram of final prices for product of assets and equivalent asset

### 3.3.4 PINN formulation for multi-dimensional case

With the benchmark in place, we now need to derive the variational inequality for the multi-dimensional case where we have  $n$  assets. This is a straightforward extension of the one-dimensional case, where the value function  $v(t, S^1, \dots, S^n)$  now takes  $n + 1$  inputs. For notation, we will express the multiple assets as a vector  $S = (S^1, \dots, S^n)^\top$ , so the value function can be written more concisely as  $v(t, S)$ . Similarly, the payoff function will be written as  $g(S^1, \dots, S^n) = g(S)$ .

The value function stays

$$v(t, S) = \sup_{\tau \in [t, T]} \mathbb{E}[e^{-r(\tau-t)} g(S_\tau) \mid S_t = S] \quad (3.55)$$

Should we choose not to exercise the option at time  $t$ , we can apply the same argument as in the one-dimensional case: considering a small  $h > 0$ , we can apply Itô's lemma on multiple dimensions to yield

$$\begin{aligned} v(t+h, S_{t+h}) = v(t, S) &+ \int_t^{t+h} \left( \frac{\partial v}{\partial t}(u, S_u) + \sum_{i=1}^n \frac{\partial v}{\partial S^i}(u, S_u) r S_u^i \right. \\ &\left. + \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \frac{\partial^2 v}{\partial S^i \partial S^j}(u, S_u) \sigma_i \sigma_j \rho_{ij} S_u^i S_u^j \right) du \end{aligned} \quad (3.56)$$

The steps from here are identical to the one-dimensional case, where we can take expectations of both sides and use the second order Taylor expansion of the exponential to get the variational inequality

$$rv(t, S) \geq \frac{\partial v}{\partial t}(t, S) + \sum_{i=1}^n \frac{\partial v}{\partial S^i}(t, S) r S^i + \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \frac{\partial^2 v}{\partial S^i \partial S^j}(t, S) \sigma_i \sigma_j \rho_{ij} S^i S^j \quad (3.57)$$

where to simplify into the PINN framework as in (2.39), we can define the operator  $\mathcal{L}$  as

$$\mathcal{L}[v] = rv(t, S) - \sum_{i=1}^n \frac{\partial v}{\partial S^i}(t, S) r S^i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \frac{\partial^2 v}{\partial S^i \partial S^j}(t, S) \sigma_i \sigma_j \rho_{ij} S^i S^j \quad (3.58)$$



and the problem reduces to the same variational inequality as in the one-dimensional case (3.37), with a different operator  $\mathcal{L}$ .

To finish the PINN formulation, we need to derive the boundary conditions. The terminal condition is straightforward, as it is given by the payoff function:

$$v(T, S) = g(S) = \left( K - \prod_{i=1}^n S^i \right)^+ \quad (3.59)$$

For the spatial boundary conditions, we will again truncate the domain to

$$\Omega = \{S = (S^1, \dots, S^n) \in \mathbb{R}^n : S^i \in (S_{\min}^i, S_{\max}^i) \text{ for } i = 1, \dots, n\} \quad (3.60)$$

for some sufficiently small and large respective minimum and maximum values for each asset. The lower boundary will be sampled from the set

$$\partial_{\text{lower}}\Omega = \{S \in \Omega : S^i = S_{\min}^i \text{ for some } i\} \quad (3.61)$$

where we have the boundary condition

$$v(t, S) = g(S), \quad S \in \partial_{\text{lower}}\Omega \quad (3.62)$$

and the upper boundary will be sampled from the set

$$\partial_{\text{upper}}\Omega = \{S \in \Omega : S^i = S_{\max}^i \text{ for some } i\} \quad (3.63)$$

with the boundary condition

$$v(t, S) = 0, \quad S \in \partial_{\text{upper}}\Omega \quad (3.64)$$

analogous to the one-dimensional case. Of course, this framework could be easily adapted to other payoff functions, where the boundary conditions would change accordingly (but not the variational inequality). This will be explored later in the numerical results when we consider a different payoff. However, for our benchmark, the put on the product of assets is the setup we will use.

### 3.4 Sobolev Deep Neural Networks

For the Black-Scholes model, we will additionally explore a different method to price American puts with neural networks, but this time instead of the traditional boundary conditions, we will work with fractional Sobolev norms to construct our loss function. This is because it has been shown that conditional on the loss converging to 0, the neural network will converge to the true solution [8], which is a stronger guarantee than just pointwise convergence. Additionally, the loss function components no longer depend on the payoff and is applicable for all situations, a desirable property for a general method.

The theory of Sobolev spaces is rich and we will not attempt to cover all the details, to avoid diverting attention from the main focus of the project. However, we make particular use of fractional Sobolev spaces, which are a generalisation of the integer order Sobolev spaces. These definitions are taken from [25, 26] which are very good resources for the underlying theory.

### 3.4.1 Sobolev spaces

Throughout this section, let  $\Omega \subset \mathbb{R}^n$  be an open bounded set with Lipschitz boundary, on which our target function is defined, and let  $p \in [1, +\infty)$  be the integrability parameter. Recall the Lebesgue space

$$L^p(\Omega) := \left\{ u : \Omega \rightarrow \mathbb{R} : u \text{ is measurable and } \int_{\Omega} |u(x)|^p dx < \infty \right\} \quad (3.65)$$

and the integer order Sobolev space  $W^{m,p}(\Omega)$  for  $m \in \mathbb{N}$  [25],

$$W^{m,p}(\Omega) := \{ u \in L^p(\Omega) : D^{\alpha}u \in L^p(\Omega), |\alpha| \leq m \}, \quad (3.66)$$

where  $D^{\alpha}u$  denotes the weak derivative of  $u$  with respect to the multi-index  $\alpha$ .

The fractional Sobolev space  $W^{s,p}(\Omega)$  is defined with the fractional exponent  $s \in (0, 1)$  as [26]:

$$W^{s,p}(\Omega) := \left\{ u \in L^p(\Omega) : \frac{|u(x) - u(y)|}{|x - y|^{n/p+s}} \in L^p(\Omega \times \Omega) \right\}, \quad (3.67)$$

which can be interpreted as an intermediate Banach space between  $L^p(\Omega)$  and  $W^{1,p}(\Omega)$ , endowed with the norm

$$\begin{aligned} \|u\|_{W^{s,p}(\Omega)}^p &= \|u\|_{L^p(\Omega)}^p + [u]_{W^{s,p}(\Omega)}^p \\ &= \int_{\Omega} |u|^p dx + \int_{\Omega} \int_{\Omega} \frac{|u(x) - u(y)|^p}{|x - y|^{n+sp}} dx dy, \end{aligned} \quad (3.68)$$

where we have used the *Gagliardo seminorm* of  $u$ :

$$[u]_{W^{s,p}(\Omega)} := \left( \int_{\Omega} \int_{\Omega} \frac{|u(x) - u(y)|^p}{|x - y|^{n+sp}} dx dy \right)^{1/p}. \quad (3.69)$$

We can extend this definition to non-integer  $s > 1$ . In the decomposition  $s = m + \sigma$  with  $m \in \mathbb{N}$  and  $\sigma \in (0, 1)$ , we define

$$W^{s,p}(\Omega) := \{ u \in W^{m,p}(\Omega) : D^{\alpha}u \in W^{\sigma,p}(\Omega), |\alpha| = m \} \quad (3.70)$$

endowed with the natural norm

$$\|u\|_{W^{s,p}(\Omega)}^p := \|u\|_{W^{m,p}(\Omega)}^p + \sum_{|\alpha|=m} [D^{\alpha}u]_{W^{\sigma,p}(\Omega)}^p. \quad (3.71)$$

In the special case  $p = 2$ , the fractional Sobolev space becomes a Hilbert space, which we denote  $H^s(\Omega) := W^{s,2}(\Omega)$ . This is the specific case we will be working with in our loss construction.

### 3.4.2 Model

Suppose we would like to price an option on  $n$  assets  $S = (S^1, \dots, S^n)$ . As asset prices are theoretically unbounded, we truncate it in our model much like the vanilla PINN. Thus the domain of the target function can be chosen to be the open set

$$\Omega = (S_{\min}^1, S_{\max}^1) \times \dots \times (S_{\min}^n, S_{\max}^n) \quad (3.72)$$

In our study we will make the simplification that  $S_{\min}^i = S_{\min}^j$  and  $S_{\max}^i = S_{\max}^j$  for all pairs  $i, j$ , which we will denote with  $a$  and  $b$  respectively. Our simplified spaces are therefore

$$\Omega = (a, b)^n \quad \Omega_T = (0, T) \times \Omega \quad \Sigma = (0, T) \times \Gamma \quad (3.73)$$

where  $\Gamma := \partial\Omega$  is the boundary of  $\Omega$ , a hyper-rectangle. We will use the representation

$$\Gamma = \bigcup_{i=1}^n (\{x \in [a, b]^n : x_i = a\} \cup \{x \in [a, b]^n : x_i = b\}) \quad (3.74)$$

which is not the simplest representation, but it is the most useful for our purposes as it is decomposed into faces, needed for the normal derivative. Note that in the one-dimensional case, the boundary  $\Gamma$  reduces simply to the endpoints  $\Gamma = \{a, b\}$ .

### 3.4.3 Loss function

Our loss function would be a weighted sum of four loss components  $\sum_{i=1}^4 \lambda_i J_i$ , each of which ensures regulation of a specific aspect of our function, with the weights either kept constant or chosen through adaptive loss weighting as presented in Section 2.3.3. We present the generalised loss construction, applicable for any dimension.

1. The variational loss: this is as before, controlling the accuracy of our function in the interior of our domain:

$$J_1 = \|\min\{\mathcal{G}[f], f - g\}\|_{L^2(\Omega_T)}^2 \quad (3.75)$$

2. The energy loss: before we were only controlling the L2 norm of the payoff, but we require the option value to equal the payoff as a function and not just pointwise. Thus we need to incorporate the gradient into the loss so that the spatial slope of the payoff is respected.

$$J_2 = \|f(T, \cdot) - g(\cdot)\|_{H^1(\Omega)}^2 \quad (3.76)$$

where we have the norm defined as

$$\|d\|_{H^1(\Omega), \mu_2}^2 = \|d\|_{L^2(\Omega)}^2 + \|\nabla d\|_{L^2(\Omega)}^2 \quad (3.77)$$

and the derivative

$$\nabla d = (\partial_{x_1} d, \dots, \partial_{x_n} d) \quad (3.78)$$

3. The Dirichlet trace loss [27]: this enforces  $f = g$  on the spatial boundary  $\Sigma$ .

$$J_3 = \|f(t, x) - g(x)\|_{H^{\frac{3}{4}, \frac{3}{2}}(\Sigma)}^2 \quad (3.79)$$

where we select the fractional Sobolev space  $H^{\frac{3}{4}, \frac{3}{2}}(\Sigma)$  from [28], and the fractional norm is defined as

$$\begin{aligned} \|d\|_{H^{\frac{3}{4}, \frac{3}{2}}(\Sigma)}^2 &= \|d\|_{H^{0,1}(\Sigma)}^2 + \int_0^T \int_0^T \frac{\|d(t_1, \cdot) - d(t_2, \cdot)\|_{L^2(\Gamma)}^2}{|t_1 - t_2|^{2.5}} dt_1 dt_2 \\ &\quad + \int_0^T \int_{\Gamma} \int_{\Gamma} \frac{|d(t, x) - d(t, y)|^2}{\|x - y\|^n} d\sigma(x) d\sigma(y) dt \end{aligned} \quad (3.80)$$

The double time integral penalises the variation of  $f$  between two time points, and the double space integral penalises the variation of  $f$  between two spatial points.

The time exponent 2.5 is equal to  $2p + 1$  where  $p$  is the decimal part of the time fractional exponent, and the space exponent is equal to  $2\theta + n - 1$  where  $\theta$  is the decimal part of the space fractional exponent, and  $n$  the dimension. Note that here  $d\sigma(x)$  and  $d\sigma(y)$  indicate the surface measures on  $\Gamma$ , but in practice we can treat these as  $dx$  and  $dy$ .

4. The Neumann trace loss [27]: this ensures the normal derivative of  $f - g$  vanishes on the boundary  $\Sigma$ .

$$J_4 = \|\nabla(f - g) \cdot \nu\|_{H^{\frac{1}{4}, \frac{1}{2}}(\Sigma)}^2 \quad (3.81)$$

where our space  $H^{\frac{1}{4}, \frac{1}{2}}(\Sigma)$  is again selected from, and  $\nu$  is the unit normal vector pointing outwards from the domain  $\Omega$  [28]. That is to say from an implementation perspective, the normal derivative is  $\nabla(f - g) \cdot \nu = \partial_{x_{i^*}}(f - g)$  with  $x_{i^*}$  the coordinate normal to the face of the boundary we are on, justifying the choice of the representation of  $\Gamma$  as a union of faces. Rigorously, we require the sign, but in practice this is not necessary, as we will take the square in the end. The fractional norm is thus written

$$\begin{aligned} \|d\|_{H^{\frac{1}{4}, \frac{1}{2}}(\Sigma)}^2 &= \|d\|_{L^2(\Sigma)}^2 + \int_0^T \int_0^T \frac{\|d(t_1, \cdot) - d(t_2, \cdot)\|_{L^2(\Gamma)}^2}{|t_1 - t_2|^{1.5}} dt_1 dt_2 \\ &\quad + \int_0^T \int_{\Gamma} \int_{\Gamma} \frac{|d(t, x) - d(t, y)|^2}{\|x - y\|^n} d\sigma(x) d\sigma(y) dt \end{aligned} \quad (3.82)$$

because  $H^{0,0} = L^2$ .

The third and fourth losses are those that enforce  $f$  to solve the option pricing variational inequality with regularity in the time and space dimensions beyond pointwise [29]. In practice, we will compute a Monte Carlo estimate of these integrals by sampling pairs  $(t_1, t_2)$  from the domain, taking care that they are not too close to each other by rejecting pairs that have a distance less than a threshold. We then evaluate the integrand and take the mean across all pairs.

## 3.5 Numerical Results


### 3.5.1 Tree benchmark against closed-form solution

First we will evaluate the performance of the benchmark binomial tree method for the one-dimensional case against the closed-form solution, which we would take as the ground truth. To do this, we will compute the tree option price without taking the maximum of the payoff and continuation value at each step so that we can compare the tree price against the Black-Scholes price.


First, we will evaluate the tree's accuracy on the most sensitive point, which is at time  $t = 0$  and when the initial asset price  $S$  is equal to the strike price  $K$  (at the money). Figure 3.2 shows the convergence of the tree price to the closed-form Black-Scholes price for European put options as we increase the depth of the tree. We can see that around 1000 steps yields a sufficiently close price to the closed-form solution, and further increasing the depth does not yield significant improvements, so this is the depth we will use for our tree benchmark.

Evaluating the tree price against the closed-form solution for different initial asset prices  $S$  at the initial time  $t = 0$  in Figure 3.3, we can see that the tree price is very close to the ground truth, with less than 0.05% of relative error until we are far out of the money where the option price is very close to 0 and unsurprisingly yields greater errors, where machine precision becomes a limiting factor.

We can see that the tree price also closely tracks the term structure of the option price, as shown in Figure 3.4, where the error is again less than 0.05%, getting even closer to zero as we get



**Figure 3.2** Tree convergence at  $t = 0, S = K$  over different depths




**Figure 3.3** Comparison of tree and closed-form prices for European put option with varying initial asset price

closer to maturity, which is expected as the tree has not done as many backward-induction steps to get to the initial time, where errors can accumulate.

Overall, we can conclude that the tree benchmark yields very strong results against the closed-form solution, giving us good confidence that it will provide accurate results for American options too, where we do not have a closed-form solution to compare against. This will serve as our benchmark for evaluating the performance of our PINN implementation for American options.

### 3.5.2 PINN results against tree benchmark


We will now train our PINN on the one-dimensional case and compare the results against the tree benchmark on the American put on the asset with parameters  $K = 1, r = 0.1, \sigma = 0.3, T = 1$  across a range of possible initial asset prices  $S_0 \in [0, 3]$ . This will be done through a feedforward neural network with 3 hidden layers each of width 32, with learning rate 0.001, batch size 10000 with 30000 maximum epochs with early stopping patience of 500 epochs. We train 25 models with different random initialisation and compute the ensemble mean for our final



**Figure 3.4** Comparison of tree and closed-form prices for European put option with varying time

price estimate.


Figure 3.5 shows the convergence of the loss function for one run, and we can see that the loss converges well, with each of them having similar magnitudes, converging in the order of  $10^{-5}$ .




**Figure 3.5** Loss convergence for the one-dimensional PINN for one run

Figure 3.6 shows the heatmap of the PINN price against the tree price across the domain of time and initial asset price. We can see that the PINN price is very close to the tree price across the domain, where the time value is correctly captured at the most sensitive point of  $S = K$  and  $t = 0$  as seen in Figure 3.7. This figure additionally shows the model capturing the term structure of the option price, where the PINN price closely tracks the tree price across time for  $S = K$ .

We can also zoom in on the most sensitive point of  $(t, S) = (0, K)$  and evaluate the error. Table 3.1 shows the mean PINN price across 30 runs with its 95% confidence interval, and we can see that the relative error is around 3%, which is good considering we were sampling across the whole domain and not just at this point. However, we can see that the benchmark tree price lies around 4 standard deviations away from the PINN mean price, suggesting that there is some bias in the PINN approximation.



**Figure 3.6** PINN vs tree price comparison heatmap across domain



**Figure 3.7** Price comparison for varying underlying price, and for term structure at  $S = K$


We can investigate this apparent bias by tracking the price  $f(t = 0, S = K)$  across different training epochs, as shown in Figure 3.8. We can see that the price spikes above the tree price and then settles around the marked benchmark tree price, approaching it from above. This suggests either that the PINN price would eventually converge to the tree price with more training, or that there is some bias in the PINN approximation that prevents it from getting closer to the tree price, which we are taking to be the ground truth.

### 3.5.3 Multi-dimensional PINN

We move onto training a PINN on multiple assets. For simplicity and visualisation purposes, we will train a two-dimensional PINN on the American put with the product of two assets as the payoff. We choose arbitrarily the parameters  $K = 1, r = 0.1, T = 1$  with asset volatilities  $\sigma = (0.3, 0.2)^\top$  and correlation between assets  $\rho = 0.5$ . Both of these assets will have a truncated interior domain  $\Omega = (0, 3)$  and we will sample uniformly from the whole domain for training. The model architecture will be similar to the one-dimensional case, with a bigger feedforward neural network of 4 hidden layers each of width 64, learning rate 0.001 and batch size 10000. This will be trained for a maximum of 40000 epochs with early stopping patience of 500 epochs across 10 different seeds. Figure 3.9 shows the heatmap of the PINN price against the tree price, where we take the  $S^1$ - $S^2$  cross section and fix time at  $t = 0$ . We can see that the PINN price is close to the tree price across the domain, with some stronger bias at the in-the-money regions where the price is higher and hence the absolute error is not as significant when we take relative error into consideration.

| Tree price | PINN price          | Relative error |
|------------|---------------------|----------------|
| 0.0833     | $0.0860 \pm 0.0006$ | 3.33%          |

**Table 3.1** PINN price at the money and initial time against tree benchmark




**Figure 3.8**  $f(0, K)$  across training epochs for one run

We can zoom in on the behaviour along the line  $S^1 \times S^2 = K$ , and plot the price along this line across time, shown in the left figure of Figure 3.7. We can see that the PINN fit is satisfactory, with inconsistencies at the tails due to our boundary conditions, which are set at  $S = 0$  and  $S = 3$  for both assets, which are not the true boundaries of the problem and a model simplification we made.

Overall, by Table 3.2, we can see that the relative error at the most sensitive point at the triple  $(0, \sqrt{K}, \sqrt{K})$  is around 7% which is considerably higher than the one-dimensional case, which is expected as we are now sampling across a much larger domain, and we are restrained by model capacity and training time due to personal hardware constraints, which is a common issue for multi-dimensional PINNs.

We can also take the cross section at  $S^1 = S^2$  to evaluate the term structure of the option price, as shown in the right figure of Figure 3.10. We can see that the PINN has managed to capture the theta decay of the option price, but the consistent bias of the price is evident here and the model struggles to capture the more rapid theta decay as we approach maturity. This may be due to the fact gamma approaches a dirac delta function at the money as we approach maturity,




**Figure 3.9** PINN vs tree price comparison heatmap across  $S^1$ - $S^2$  cross section at  $t = 0$



| Tree price | PINN price          | Relative error |
|------------|---------------------|----------------|
| 0.0993     | $0.1060 \pm 0.0048$ | 6.75%          |

**Table 3.2** Prices  $f(0, \sqrt{K}, \sqrt{K})$  for the two-dimensional PINN against tree benchmark

whose sharp feature is difficult for the model to capture when our samples are sparser in the expanded domain.




**Figure 3.10** PINN vs tree price comparison for at the money cross section and term structure

### 3.5.4 Sobolev PINN results


#### One-dimensional case

We will now train a PINN with the Sobolev loss function, with the same model architecture and training setup as the vanilla PINN, and compare the results against the tree benchmark. We can see from Figure 3.11 that the Sobolev PINN price also gets close to the tree price across the domain, but the inaccuracy here is more significant than the vanilla PINN, especially at the  $(t = 0, S = K)$  point, which is arguably the point we are most interested in. However, considering that the boundary conditions of the loss function isn't manufactured from the payoff function directly, we can see that this method could be effective for more general problems where boundary conditions are not as straightforward to implement, and the price is still in the right ballpark.



**Figure 3.11** Sobolev PINN vs tree price comparison heatmap across domain

We can see especially in Figure 3.12 that the Sobolev PINN manages to capture the term structure of the option price with the theta collapse as we approach maturity, and the error is more significant due to a hump in the price against asset price around the strike price. This suggests that the model is capable of a good fit across the domain, but may be undertrained or have too small a capacity.




**Figure 3.12** Price comparison for varying underlying price, and for term structure at  $S = K$  with the Sobolev PINN

Overall this method is promising, but it may require more careful training and tuning to get better results as presented in [8], who demonstrated its effectiveness through extensive training on the GPU with much larger networks, which is beyond our scope.

### Multi-dimensional case

We will finish this section by presenting our results for the multi-dimensional case, where we have added a L-BFGS optimiser after having trained with Adam for a good number of epochs.

Figure 3.13 shows the heatmap of the Sobolev PINN price against the tree price across the  $S^1$ - $S^2$  cross section at  $t = 0$ . We can see that the Sobolev PINN price has captured the correct shape of the price across the domain, with greater inaccuracy than the vanilla PINN.




**Figure 3.13** Sobolev PINN vs tree price comparison heatmap across  $S^1$ - $S^2$  cross section at  $t = 0$

We can see that the bias is the strongest on the at-the-money line, which we can examine more closely in the left figure of Figure 3.14. We see a similar behaviour as the vanilla case exhibited

here, where the PINN overprices for more moderate values of  $S^1$  and underprices for more extreme values of  $S^1$ . This brings the case of training an enlarged domain into question, where we only evaluate our model on a smaller interior of the domain.

The right figure of Figure 3.14 shows the term structure of the option price, with the asset prices fixed at  $S^1 = S^2 = \sqrt{K}$ . We can see that the Sobolev PINN captures the theta decay of the option price better than the vanilla PINN, but there is a systematic overpricing bias across the whole time dimension. We have attempted to combat this by noticing that the errors cumulate backwards from expiry  $t = T$  backwards to the initial time  $t = 0$ , so we have attempted to train the model with causal weighting instead of a uniform average, where we put more weight on the later time steps until the residual loss is sufficiently small, and then relax this causal weighting gradually to a uniform average. This has improved the accuracy from over 100% relative error to around 20% relative error, which is not ideal but is still a significant improvement nonetheless.



**Figure 3.14** Price comparison for varying underlying price, and for term structure at  $S = K$  with the Sobolev PINN

Even though this approach yielded lower accuracy than that of the vanilla PINN, we could still see its strengths over the vanilla PINN (such as it being capable of capturing the collapse of the option price to zero as we approach maturity). With more careful training and tuning, the accuracy could definitely improved, which is an avenue for future work. We will prefer the vanilla PINN for the free boundary calculation due to its better accuracy, but it was interesting nonetheless to see how we could construct the loss function from another angle altogether, without relying on manual boundary conditions.


### 3.5.5 Free boundary

#### One-dimensional case

For the one-dimensional free boundary, we will first evaluate the free boundary from the tree benchmark. Here, the continuation values are extracted through a small modification to the tree algorithm, where we compute the price of the Bermudan option with the same parameters with possibility of exercise at every time step except for at  $t = 0$ , effectively serving as a continuation value of the American option at  $t = 0$ . This continuation value is compared with the intrinsic value (the payoff) like we introduced in equation (2.12) in the mathematical background. The left figure in Figure 3.15 shows the continuation probabilities across different  $(t, S)$  pairs, and here we use hard classification as binomial tree pricing is deterministic.


We can see that the free boundary is correctly captured at glance, where we exercise the option when the asset price is sufficiently low, whose threshold increases as we get closer to maturity, consistent with the intuition that the option is losing time value.

We now evaluate the free boundary from the PINN. The central figure in Figure 3.15 shows the continuation probabilities across different  $(t, S)$  pairs using the sigmoid distribution as described in equation (2.15). We can see that the shape of the free boundary is similar to the tree benchmark, and with soft classification the lowest probability is around 0.4, which is not unintuitive as in the exercise region, the continuation value is theoretically equal to the payoff, and hence we should be indifferent between exercising and continuing.



**Figure 3.15** Estimated continuation probabilities from PINN for the American put option

To use the Gaussian method, we will first evaluate its validity by plotting the distribution of the standardised price estimates across all seeds in Figure 3.16, noting that as the variance is unknown and we use the sample variance, the standardised price estimates would follow a t-distribution with  $n - 1$  degrees of freedom where  $n$  is the number of seeds.



**Figure 3.16** Distribution of standardised neural network price estimates across 25 runs


By looking at the histogram and the Q-Q plot, we can see that the distribution of the standardised price estimates is reasonably close to a t-distribution, with the exception of the tails. We see that the observed distribution has lighter tails than the theoretical t-distribution, which is more statistically acceptable, so we will nevertheless proceed with caution in evaluating the free boundary using the Gaussian method. The right figure in Figure 3.15 shows the continuation probabilities across different  $(t, S)$  pairs.

We can see that the shape roughly corresponds to the free boundaries we have seen before, but this is a much noisier estimate, which puts into question the assumptions we have made.

We are essentially assuming that the price estimates of the PINN are unbiased and normally distributed around the true price, which is quite strong and not true in practice, as we have already observed in Figure 3.6 where we have zones where the PINN price is systematically above or below the tree price, suggesting bias.


## Two-dimensional case

Using the same methodology as the one-dimensional case, we plot our estimated free boundaries in Figure 3.17. This is done on the  $S^1$ - $S^2$  cross section at  $t = 0$ , with the additional reference line of  $S^1 \times S^2 = K$ , which should be a continuation region as it is at the money and is worthless under immediate exercise.



**Figure 3.17** Estimated continuation probabilities for 2-dimensional case with  $S^1 \times S^2 = K$  reference line

We can see that both methods of estimating the continuation probabilities yield similar shapes to the binomial tree benchmark, with the Gaussian method being noisier than the sigmoid method. We can further examine the continuation probabilities over time along the  $S^1 = S^2$  cross section, varying across time this time. This is shown in Figure 3.18, where we can see that the PINN has successfully captured the structure of the free boundary across time where  $S^1 = S^2$ .



**Figure 3.18** Estimated continuation probabilities for 2-dimensional case on  $S^1 = S^2$  cross section across time

## 3.6 Minimum of two assets

To demonstrate the flexibility of our PINN framework, we will now shift our focus onto different payoffs on multiple assets as an extension. [30] explores a variety of different payoffs for

American options on multiple assets, but that with a particularly interesting free boundary is the one of the call on the maximum of two assets. We will deal with the analogous case for the put on the minimum of two assets.

To price this, we can simply change the payoff function in our loss function, and update the boundary conditions accordingly to the new payoff function, but the variational inequality and the rest of the framework remain unchanged. For the American put on the minimum of two assets, the payoff function is given by  $g(S^1, S^2) = (K - \min(S^1, S^2))^+$ , and thus we must have the terminal boundary condition

$$v(T, S^1, S^2) = (K - \min(S^1, S^2))^+ \quad (3.83)$$

When one of the assets is very high (i.e.,  $S^1 \rightarrow \infty$  or  $S^2 \rightarrow \infty$ ), the minimum of the two assets is just the other asset, and thus we have


$$v(t, \infty, S^2) = v_{1D}(t, S^2) \quad v(t, S^1, \infty) = v_{1D}(t, S^1) \quad (3.84)$$

where  $v_{1D}$  is the price of the one-dimensional American put option on a single asset, which we can compute with the different methods we have already explored. Finally, when one of the assets is worthless, the minimum of the two assets is just the worthless asset, and hence the value will be  $K$  as we should exercise immediately. Thus, the boundary condition would be

$$v(t, S^1, S^2) = K \quad \text{for } S^1 = 0 \text{ or } S^2 = 0 \quad (3.85)$$

Together, these boundary conditions (alongside the variational inequality which remains the same with the operator  $\mathcal{L}$  as in 3.58) form the PINN framework which we can train.


In Figure 3.19, we can see that our PINN yields results that are in line with our expectations. We expect to see some symmetry along the  $S^1 = S^2$  line as the payoff is symmetric in the two assets. We introduce some asymmetry by training the network where  $S^2$  has a higher constant volatility than  $S^1$ .



**Figure 3.19** Heatmap for the put on the minimum of two assets across the  $S^1$ - $S^2$  cross section at  $t = 0$ , and the  $S^1$ - $t$  cross section with  $S^1 = S^2$

We can also see that the term structure is somewhat captured by the contour lines of the right figure of Figure 3.19, where the price decays as we approach maturity. We train the model with a high volatility of  $(\sigma_1, \sigma_2) = (0.4, 0.5)$  to make the theta decay more significant. This also gives us a clearer evolution of the free boundary, which we see in Figure 3.20, where we can see that not only the continuation region is a mirror image of the theoretical result described in

[30] as we would expect, but the PINN also captures the shrinking nature of the continuation region as we approach maturity, consistent with the intuition that the option loses time value and the region would close entirely, leaving only the box  $\{(S^1, S^2) : S^1 \geq K, S^2 \geq K\}$  as the continuation region at maturity.



**Figure 3.20** Continuation probabilities for the put on the minimum of two assets across the  $S^1$ - $S^2$  cross section with varying  $t$

Overall, this demonstrates the capacity of our PINN approach to pricing options and we can see that the framework is flexible enough to be applied to different payoffs with minimal adjustments, yielding convincing results that are consistent with theoretical results. For the next step, we will move onto a different model, the Heston model, which injects randomness into the volatility, which is a more realistic assumption.

## Chapter 4

# The Heston Model

The classic Black Scholes model makes the strong assumption that the stock returns are normally distributed with known mean and variance. However, the Black Scholes formula does not depend on the mean of the log returns  $\mu$  (since we perform a change of measure, and it depends instead on the risk-free rate  $r$ ). We therefore have two directions in which we can take this project: we can model the risk-free rate as a stochastic process, or we can allow the volatility to vary according to a stochastic process. We will focus on the latter. Heston [2] proposes a model involving a stochastic variance process in which a closed-form solution for European call options exists. In this model, the dynamics under the pricing measure of the asset price  $S$  and the variance (squared volatility) process  $V$  are governed by the stochastic differential equation system under the risk-neutral measure  $\mathbb{Q}$ :

$$dS_t = rS_t dt + S_t \sqrt{V_t} dW_t^1 \quad (4.1)$$

$$dV_t = \kappa(\theta - V_t)dt + \sigma \sqrt{V_t} dW_t^2 \quad (4.2)$$

where we have initial conditions  $S_0 = s > 0$  and  $V_0 = v \geq 0$ , and the two Brownian motions are correlated with

$$d\langle W^1, W^2 \rangle_t = \rho dt \quad \rho \in (-1, 1) \quad (4.3)$$

The dynamics of the variance  $V$  follow a Cox-Ingersoll-Ross (CIR) process with mean reversion rate  $\kappa > 0$ , long run state  $\theta > 0$  and volatility of the volatility  $\sigma > 0$ . We can enforce the Feller condition  $2\kappa\theta \geq \sigma^2$  so that the volatility process stays positive.


The Heston model takes into account the stochastic nature of volatility while enforcing its mean-reverting nature, which is observed empirically in the market. It supports non-constant volatility smiles, as we can observe in Figure 4.1. The curvature is determined by the volatility of volatility  $\sigma$ , and the skew is determined by the correlation  $\rho$ .

### 4.1 Closed-form Solution

There exists a closed-form solution derived by Heston [2] to the price of the European call option (and indeed the put option, by put-call parity). This will serve as a benchmark for our other techniques, which would give us confidence in the accuracy of our other techniques for European options, and by extension, for American options too.

The closed-form solution is derived by using an Ansatz inspired by the Black-Scholes formula, where the unknown functions in the Ansatz can be determined by substituting our Ansatz into the corresponding PDE and boundary conditions. We will include a sketch of the derivation





**Figure 4.1** Volatility smile generated by the Heston model

here which follows those found in [2, 31], but with some further justification for a better overall understanding.

Let  $C$  denote the price of the European call option. We will derive this PDE when we derive the variational equation in Section 4.3.1, where we show that the price of the European call option with strike  $K$  and maturity  $T$  must satisfy

$$\frac{\partial C}{\partial t} + rS \frac{\partial C}{\partial S} + \kappa(\theta - V) \frac{\partial C}{\partial V} + \frac{1}{2} V S^2 \frac{\partial^2 C}{\partial S^2} + \rho \sigma V S \frac{\partial^2 C}{\partial S \partial V} + \frac{1}{2} \sigma^2 V \frac{\partial^2 C}{\partial V^2} - rC = 0 \quad (4.4)$$

with the terminal condition  $C(T, S, V) = (S - K)^+$ .

We will mimic the form of the Black-Scholes formula to form the Ansatz for the solution

$$C(S, V, t) = S P_1(x, V, t) - K e^{-r(T-t)} P_2(x, V, t), \quad x := \ln S \quad (4.5)$$

with terminal conditions  $P_j(x, V, T) = \mathbf{1}_{\{x \geq \ln K\}}$  so that our Ansatz reduces to  $(S - K)^+$  at maturity. We can substitute the Ansatz into the PDE and equate the coefficients of  $S$  and  $K e^{-r(T-t)}$  to obtain two PDEs for  $P_1$  and  $P_2$ :

$$\frac{\partial P_j}{\partial t} + (r + u_j V) \frac{\partial P_j}{\partial x} + (a - b_j V) \frac{\partial P_j}{\partial V} + \frac{1}{2} V \frac{\partial^2 P_j}{\partial x^2} + \rho \sigma V \frac{\partial^2 P_j}{\partial x \partial V} + \frac{1}{2} \sigma^2 V \frac{\partial^2 P_j}{\partial V^2} = 0 \quad (4.6)$$

where  $u_1 = \frac{1}{2}$ ,  $u_2 = -\frac{1}{2}$ ,  $a = \kappa\theta$ ,  $b_1 = \kappa - \rho\sigma$  and  $b_2 = \kappa$  [2]. With this in mind, let's define the characteristic functions

$$\phi_j(x, V, t; \omega) = \mathbb{E}^{\mathbb{Q}_j} [e^{i\omega x_T} | \mathcal{F}_t] \quad (4.7)$$

which admits the Ansatz

$$\phi_j(x, V, t; \omega) = \exp(C_j(\tau, \omega) + D_j(\tau, \omega)V + i\omega x), \quad \tau := T - t \quad (4.8)$$

since the coefficients in the PDE for  $P_j$  are independent of  $x$  (justifying the  $e^{i\omega x}$  term) and affine in  $V$  (justifying the  $e^{D_j V}$  term) [32].

We can insert this Ansatz into the PDE for  $P_j$  to get a system of ODEs for  $C_j$  and  $D_j$ :

$$\frac{\partial D_j}{\partial \tau} = \frac{1}{2} \sigma^2 D_j^2 + (\rho \sigma i \omega - b_j) D_j + u_j i \omega - \frac{1}{2} \omega^2, \quad \frac{\partial C_j}{\partial \tau} = r i \omega + a D_j \quad (4.9)$$

with the boundary conditions  $C_j(0, \omega) = D_j(0, \omega) = 0$ . There exists a closed-form solution to this system of ODEs, where we first define

$$d_j := \sqrt{(\rho\sigma i\omega - b_j)^2 - \sigma^2(2u_j i\omega - \omega^2)} \quad (4.10)$$

$$g_j := \frac{b_j - \rho\sigma i\omega - d_j}{b_j - \rho\sigma i\omega + d_j} \quad (4.11)$$

where we use the opposite sign convention compared to Heston's original formulation for better numerical stability (more precisely to avoid regions of numerical discontinuity in the complex plane) as suggested in [33]. This yields the solutions

$$D_j(\tau, \omega) = \frac{b_j - \rho\sigma i\omega - d_j}{\sigma^2} \cdot \frac{1 - e^{-d_j\tau}}{1 - g_j e^{-d_j\tau}} \quad (4.12)$$

$$C_j(\tau, \omega) = r i\omega\tau + \frac{a}{\sigma^2} \left[ (b_j - \rho\sigma i\omega - d_j)\tau - 2 \ln \frac{1 - g_j e^{-d_j\tau}}{1 - g_j} \right] \quad (4.13)$$

By the Gil-Pelaez inversion theorem [34], we can retrieve  $P_j$  from its characteristic function  $\phi_j$ :

$$P_j = \frac{1}{2} + \frac{1}{\pi} \int_0^\infty \Re \left[ \frac{e^{-i\omega \ln K} \phi_j(x, V, t; \omega)}{i\omega} \right] d\omega \quad (4.14)$$

which we can substitute back into our original Ansatz for the call option (4.5) to get the closed-form solution for the call option price. If we would like to be pedantic, we could argue this is not strictly closed form as we require numerical integration to compute the second term of (4.14), but this remains a semi-analytical and tractable solution nonetheless, that we could use as a benchmark for our other methods.

## 4.2 Tree Solution

Since the closed-form solution is limited to European options, we would like to find ways to approximate solutions to the pricing problem so that we can eventually extend the case to American options, which has no closed-form solution. The first method we will explore is the tree method, where it is easy to modify to the American case by adding an additional comparison to the intrinsic value at each node. We will use the closed-form formula as a benchmark for the tree method for European options, to give us confidence in the solution the tree model provides for American options as it is a simple modification in the algorithm, which we can use as a benchmark for the neural networks we will explore later.

The idea will be similar to the binomial tree we explored earlier in the paper for the Black-Scholes model, but we now have two correlated stochastic processes in the Heston model: the asset price process and the variance process. As this tree is not recombining, generating paths for these processes and computing the option prices by backwards induction naively would lead to an exponential time complexity, which is not computationally feasible. Therefore, we will use a grid interpolation approach as in [4] so that the number of nodes grows quadratically in the number of time steps instead.

For the algorithm, we will work under the variance process  $V$  and log stock price process  $Z$  (which is a simple transformation from  $S$  using Itô's formula):

$$dV_t = \kappa(\theta - V_t)dt + \sigma\sqrt{V_t}dW_t^1 \quad (4.15)$$

$$dZ_t = (r - \frac{1}{2}V_t)dt + \sqrt{V_t}dW_t^2 \quad (4.16)$$

where  $d\langle W^1, W^2 \rangle_t = \rho dt$  under the risk-neutral measure  $\mathbb{Q}$  and  $\kappa, \sigma, \theta, r > 0$  as before. Of course, the value  $C$  of the European option with maturity  $T$  and payoff function in terms of stock and volatility values  $g : \mathbb{R}^+ \times \mathbb{R}^+ \rightarrow \mathbb{R}$  will be

$$C(t, V_t, Z_t) = e^{-r(T-t)} \mathbb{E}^{\mathbb{Q}}[g(e^{Z_t}, \sqrt{V_t}) \mid \mathcal{F}_t] \quad (4.17)$$

#### 4.2.1 Euler scheme

To construct the tree, we need to discretise the process. Suppose we have  $n$  steps, then  $\Delta t = \frac{T}{n}$ . Using the standard Euler scheme, we can have

$$V_{k+1} = V_k + \kappa(\theta - V_k^+) \Delta t + Y_{k+1}^1 \sigma \sqrt{V_k^+ \Delta t} \quad (4.18)$$

$$Z_{k+1} = Z_k + (r - \frac{1}{2} V_k) \Delta t + Y_{k+1}^2 \sqrt{V_k^+ \Delta t} \quad (4.19)$$

where the increment variables  $(Y_k^1, Y_k^2)$  are iid pairs across time steps  $k$  with the joint distribution given by

$$\mathbb{Q}^n(Y_k^1 = i_1, Y_k^2 = i_2) = \frac{1}{4}(1 + i_1 i_2 \rho) \quad i_1, i_2 \in \{-1, 1\} \quad (4.20)$$

We use a different measure  $\mathbb{Q}^n$  (with  $n$  indicating the number of time steps) as it is not the same as the risk-neutral measure in the Heston model. To ensure positivity of the process, we have taken the positive part of the variance process  $V_k^+ := \max(0, V_k)$ . Note here,  $V_k$  shouldn't be negative anyway since we will enforce the Feller condition in practice, but with this in place we have full generality.

#### 4.2.2 Interpolation optimisation

The problem with the naïve Euler scheme is that the tree is not recombining, and hence the number of nodes scales exponentially with increasing time steps, which is not feasible computationally. We will hence, for each time step  $k$ , construct a grid of option prices at time  $k$  and the backward induction would be a bilinear interpolation of the four nearest points.

To set this up, we need the boundaries of the grid at each time step  $k$ , which is given by the minimum and maximum values of  $V_k$  and  $Z_k$  in the support. In other words, we have

$$z_k^{\max} = \max\{z : \mathbb{Q}^n(Z_k = z) > 0\} \quad (4.21)$$

$$z_k^{\min} = \min\{z : \mathbb{Q}^n(Z_k = z) > 0\} \quad (4.22)$$

$$v_k^{\max} = \max\{v : \mathbb{Q}^n(V_k = v) > 0\} \quad (4.23)$$

$$v_k^{\min} = \min\{v : \mathbb{Q}^n(V_k = v) > 0\} \quad (4.24)$$

With these boundaries, we have the grid spacing given our mesh sizes  $m_z, m_v \in \mathbb{N}^+$ :

$$\Delta z_k = (z_k^{\max} - z_k^{\min}) / m_z \quad (4.25)$$

$$\Delta v_k = (v_k^{\max} - v_k^{\min}) / m_v \quad (4.26)$$

and the gridpoints:

$$\hat{S}_k = \left\{ \left( v_k^{\min} + i \Delta v_k, z_k^{\min} + j \Delta z_k \right) \mid i \in \{0, \dots, m_v\}, j \in \{0, \dots, m_z\} \right\} \quad (4.27)$$

We will now define the discretised process  $(\tilde{V}, \tilde{Z})$  on the grid space  $\hat{\mathcal{S}}_k$ . This will involve the Euler scheme functions

$$\nu(y, v, z) = v + \kappa(\theta - v^+) \Delta t + y \sigma \sqrt{v^+ \Delta t} \quad (4.28)$$


$$\zeta(y, v, z) = z + (r - \frac{1}{2}v) \Delta t + y \sqrt{v^+ \Delta t} \quad (4.29)$$

where we notice  $\nu$  doesn't depend on  $z$ , but I will put  $z$  in as an argument nonetheless for uniform function signature.

Then using these functions, we can retrieve the next step. We can evolve our process to the next step with

$$(\tilde{V}_{k+1}, \tilde{Z}_{k+1}) = \Phi_{J_{k+1}^v, J_{k+1}^z}^{\hat{\mathcal{S}}_{k+1}}(\nu(Y_{k+1}^1, \tilde{V}_k, \tilde{Z}_k), \zeta(Y_{k+1}^2, \tilde{V}_k, \tilde{Z}_k)) \quad (4.30)$$

Where  $\Phi_{j_1, j_2}^{\mathcal{S}} : \mathbb{R}^2 \rightarrow \mathcal{S} \subset \mathbb{R}^2$  is a function mapping  $(v, z) \in \mathbb{R}^2$  to one of the grid points in  $\mathcal{S} := \mathcal{S}_v \times \mathcal{S}_z$ .  $j_1, j_2 \in \{0, 1\}$  are additional parameters that encode the snapping direction and determine the corner in which  $(v, z)$  is mapped, corresponding to  $J_{k+1}^v, J_{k+1}^z$  in (4.30). Visually, this is illustrated in Figure 4.2, where each point in the price-variance space can evolve to four different points in the next time step, each of which we take the interpolation of the four nearest grid points.



**Figure 4.2** Interpolation illustration (adapted from [4])

To do this, we define the 1D snapping functions  $\phi_i^{\mathcal{T}} : \mathbb{R} \rightarrow \mathcal{T}$  where we have the one-dimensional grid points  $\mathcal{T} \subset \mathbb{R}$  and the snap-right indicator  $i \in \{0, 1\}$ . The function is defined mathematically as

$$\phi_0^{\mathcal{T}}(u) = \max\{u' \in \mathcal{T} : u' \leq u\} \quad (4.31)$$

$$\phi_1^{\mathcal{T}}(u) = \min\{u' \in \mathcal{T} : u' \geq u\} \quad (4.32)$$

Our 2D snapping function  $\Phi_{j_1, j_2}^{\mathcal{S}} : \mathbb{R}^2 \rightarrow \mathcal{S}$  where  $\mathcal{S} = \mathcal{S}_v \times \mathcal{S}_z \subset \mathbb{R}^2$  is thus

$$(v, z) \mapsto (\phi_{j_1}^{\mathcal{S}_v}(v), \phi_{j_2}^{\mathcal{S}_z}(z)) \quad (4.33)$$

The joint probabilities [4] for  $Y_{k+1}$  are given by, for  $i_1, i_2 \in \{-1, 1\}, j_1, j_2 \in \{0, 1\}$ :

$$q_{i_1, i_2, j_1, j_2}(\tilde{V}_k, \tilde{Z}_k) := \tilde{\mathbb{Q}}(Y_{k+1}^1 = i_1, Y_{k+1}^2 = i_2, J_{k+1}^v = j_1, J_{k+1}^z = j_2 \mid (\tilde{V}_k, \tilde{Z}_k)) \quad (4.34)$$

$$= \frac{1}{4} (1 + i_1 i_2 \rho) \times c_{j_1, j_2}^{\hat{\mathcal{S}}_{k+1}} \left( \nu(i_1, \tilde{V}_k, \tilde{Z}_k), \zeta(i_2, \tilde{V}_k, \tilde{Z}_k) \right) \quad (4.35)$$

where we have defined the measure  $\tilde{\mathbb{Q}}$  as an extension of  $\mathbb{Q}^n$  to include the random snapping indicators  $J_{k+1}^v, J_{k+1}^z$  as well, with the snapping probabilities given by the interpolation quantity [35] on a general grid  $\mathcal{S} = \mathcal{S}_v \times \mathcal{S}_z \subset \mathbb{R}^2$ :

$$c_{j_1, j_2}^{\mathcal{S}}(v, z) = \tilde{v}^{j_1} (1 - \tilde{v})^{1-j_1} \cdot \tilde{z}^{j_2} (1 - \tilde{z})^{1-j_2} \quad (4.36)$$

where we have the difference terms

$$\tilde{v} = \frac{v - \phi_0^{\mathcal{S}_v}(v)}{\phi_1^{\mathcal{S}_v}(v) - \phi_0^{\mathcal{S}_v}(v)} \quad (4.37)$$

$$\tilde{z} = \frac{z - \phi_0^{\mathcal{S}_z}(z)}{\phi_1^{\mathcal{S}_z}(z) - \phi_0^{\mathcal{S}_z}(z)} \quad (4.38)$$

By using the fact our grid is evenly spaced at all time steps  $k$ , this simplifies nicely to

$$\tilde{v} = \frac{v - \phi_0^{\mathcal{S}_v}(v)}{\Delta v_k}, \quad \tilde{z} = \frac{z - \phi_0^{\mathcal{S}_z}(z)}{\Delta z_k} \quad (4.39)$$

because our grid is evenly spaced at all time steps  $k$ .

### 4.2.3 Backwards induction

We can now work backwards on the tree, having constructed the grids. Our base case is the final grid, where the price is just the payoff as we are at maturity. For a general time  $t_k$  at step  $k$  and node  $(v, z) \in \hat{\mathcal{S}}_k$ , we will generate the four possible steps determined by  $i_1, i_2$  each taking values in  $\{-1, 1\}$ , and evolve our node with

$$(v, z) \mapsto (\nu(i_1, v, z), \zeta(i_2, v, z)) =: (v'_{i_1}, z'_{i_2}) \quad (4.40)$$

Which will be in between four points on the grid  $\hat{\mathcal{S}}_{k+1}$ . These can be found by performing the  $\Phi$  function on our evolved node: for  $j_1, j_2 \in \{0, 1\}$  define the multi-index  $\alpha = (i_1, i_2, j_1, j_2)$  and we have

$$(v_\alpha, z_\alpha) = \Phi_{j_1, j_2}^{\hat{\mathcal{S}}_{k+1}}(v'_{i_1}, z'_{i_2}) \quad (4.41)$$

Each of these will have a known option price. Therefore, we can just compute the discounted expectation of these sixteen nodes (two choices for each element of  $\alpha$ ) for the option price at the original node:

$$C(t_k, v, z) = e^{-r\Delta t} \sum_{\alpha \in \mathcal{A}} q_\alpha(v, z) \times C(t_{k+1}, v_\alpha, z_\alpha) \quad (4.42)$$

where  $\mathcal{A} = \{-1, 1\}^2 \times \{0, 1\}^2$ .

### 4.2.4 Implementation techniques

If we implemented this method naively, it would be very difficult for us to use the binomial tree as a benchmark because it would need to construct a new tree from scratch for every single value of  $S_0, V_0$  and  $T$ . Thus any heatmaps or plots we create would require hundreds of trees which is not feasible, considering each tree requires many time steps and large mesh sizes to be accurate.

Thus in addition to vectorising our code for each tree, we will instead consider an initial range  $[S_0^{\min}, S_0^{\max}]$  and  $[V_0^{\min}, V_0^{\max}]$  which we will use to construct our grid bounds. More specifically,

$z_k^{\max}$  and  $v_k^{\max}$  will be the up-most path starting with  $(S_0^{\max}, V_0^{\max})$  instead of simply some individual value  $(S_0, V_0)$ , and similar for the down paths. In this way, when we use the binomial tree as a benchmark, we can use any initial value of  $(S_0, V_0)$  as long as it's in the original range and compute the price with an interpolation as described above. With this method, we can evaluate across different initial values while constructing just one tree. We sacrifice a little accuracy in doing so because the grids would be more widely spaced, which we can combat by increasing the mesh sizes. Thus, we effectively construct one more complex tree instead of hundreds of simpler ones, which is much more computationally efficient.

As for selecting our mesh sizes, we will use the result that if we have

$$\lim_{n \rightarrow \infty} \Delta t = 0, \quad \lim_{n \rightarrow \infty} \max_{k=1, \dots, n-1} \frac{\Delta v_k^n}{\Delta t_k} = 0, \quad \lim_{n \rightarrow \infty} \max_{k=1, \dots, n-1} \frac{\Delta z_k^n}{\Delta t_k} = 0 \quad (4.43)$$

then the process represented in our Heston tree converges weakly to the original process  $(V, Z)$  [4]. Note here we denote  $\Delta v_k$  from before as  $\Delta v_k^n$  to emphasize we have split our maturity  $T$  into  $n$  steps, and similarly for  $\Delta z_k^n$ .

Bearing this in mind, a resolution of  $n = 100, m_v = 300, m_z = 600$  allows our tree to faithfully reproduce the closed-form solution, which are presented in the numerical results in Section 4.5.1.

## 4.3 PINN Solution

### 4.3.1 Variational equation

We would now like to derive the variational equation for options on assets under the Heston model. Note that for a  $d$ -dimensional general process  $X_t \in \mathbb{R}^d$  solving the Itô SDE

$$dX_t = b(X_t)dt + \sigma(X_t)dW_t \quad (4.44)$$

where we have the drift  $b : \mathbb{R}^d \rightarrow \mathbb{R}^d$ , the volatility  $\sigma : \mathbb{R}^d \rightarrow \mathbb{R}^{d \times m}$  and the  $m$ -dimensional Brownian motion  $W_t$ . On second-order smooth functions  $f \in \mathcal{C}^2(\mathbb{R}^d)$ , the infinitesimal generator  $\mathcal{L}$  is defined as [36]

$$\mathcal{L}[f](x) := \lim_{t \rightarrow 0} \frac{\mathbb{E}[f(X_t) \mid X_0 = x] - f(x)}{t} \quad (4.45)$$

which exists and can be derived with Itô's formula (or Taylor's expansion) to be [12]

$$\mathcal{L}[f](x) = \sum_{i=1}^d b_i \frac{\partial f}{\partial x_i} + \frac{1}{2} \sum_{i=1}^d \sum_{j=1}^d \Sigma_{ij} \frac{\partial^2 f}{\partial x_i \partial x_j} \quad (4.46)$$

where  $b_i$  is the  $i$ -th component of the drift vector  $\mathbf{b} := b(x)$  and  $\Sigma_{ij}$  is the  $(i, j)$ -th component of the covariance matrix  $\Sigma(x) = \sigma \sigma^\top$ , with  $\sigma := \sigma(x)$ .

Let's apply this to our Heston model. The drift vector is

$$b(S, V) = \begin{pmatrix} rS \\ \kappa(\theta - V) \end{pmatrix} \quad (4.47)$$

To compute  $\sigma \sigma^\top$ , we need to represent our correlated Brownian motions in terms of independent ones. Suppose we have independent Brownian motions  $B^1$  and  $B^2$ , we can define

$$W^1 = B^1 \quad (4.48)$$

$$W^2 = \rho B^1 + \sqrt{1 - \rho^2} B^2 \quad (4.49)$$

which has the quadratic variation  $d\langle W^1, W^2 \rangle_t = \rho dt$  as required. Using these new Brownian motions, we can rewrite the SDE as

$$d \begin{pmatrix} S_t \\ V_t \end{pmatrix} = \underbrace{\begin{pmatrix} rS_t \\ \kappa(\theta - V_t) \end{pmatrix}}_{b(S_t, V_t)} dt + \underbrace{\begin{pmatrix} S_t \sqrt{V_t} & 0 \\ \rho \sigma \sqrt{V_t} & \sigma \sqrt{V_t} \sqrt{1 - \rho^2} \end{pmatrix}}_{\sigma(S_t, V_t)} \begin{pmatrix} dB_t^1 \\ dB_t^2 \end{pmatrix} \quad (4.50)$$

which means that our covariance matrix is now

$$\sigma = \sigma(S, V) = \begin{pmatrix} S\sqrt{V} & 0 \\ \rho\sigma\sqrt{V} & \sigma\sqrt{V}\sqrt{1 - \rho^2} \end{pmatrix}. \quad (4.51)$$

which means

$$\Sigma = \sigma \sigma^\top = \begin{pmatrix} S^2 V & \rho \sigma S V \\ \rho \sigma S V & \sigma^2 V \end{pmatrix} \quad (4.52)$$

Note here we have many forms of  $\sigma$  which may be confusing:  $\sigma$  without any arguments is the volatility of the variance process,  $\sigma(S, V)$  is the volatility matrix of the SDE, and  $\sigma$  is the same as  $\sigma(S, V)$  but with a bold symbol to distinguish it from the scalar  $\sigma$ .

We now have everything we need for the formula. For any function  $f = f(S, V) \in \mathcal{C}^2(\mathbb{R}^d)$ , we have

$$\mathcal{L}[f] = rS \frac{\partial f}{\partial S} + \kappa(\theta - V_t) \frac{\partial f}{\partial V} + \frac{1}{2} \left( S^2 V \frac{\partial^2 f}{\partial S^2} + 2\rho\sigma S V \frac{\partial^2 f}{\partial S \partial V} + \sigma^2 V \frac{\partial^2 f}{\partial V^2} \right) \quad (4.53)$$

This is our infinitesimal generator  $\mathcal{L}$  acting on smooth functions  $f$ . Suppose we have a European option with payoff  $g : \mathbb{R}^+ \rightarrow \mathbb{R}^+$ , then our price function  $u(t, S, V)$ , the value of the call option with expiry  $T$  at time  $t$  with  $S_t = S$  and  $V_t = V$  satisfies

$$\mathcal{G}[u] := ru - \frac{\partial u}{\partial t} - \mathcal{L}u = 0 \quad (4.54)$$

with terminal condition  $u(T, S, V) = g(S)$ . This matches the PDE we had stated in the closed-form solution in (4.4).

As a check, we can set  $\kappa = 0$  and  $\sigma = 0$  so that  $dV = 0$  and volatility becomes a constant, and our model collapses to what we have been working with before. Our generator becomes

$$\mathcal{L}[f] = rS \frac{\partial f}{\partial S} + \frac{1}{2} S^2 V_0 \frac{\partial^2 f}{\partial S^2} \quad (4.55)$$

which matches our expression (3.35) from the Black-Scholes section, since  $V_0$  is the initial (and hence constant) variance.

We have derived the variational equation for the European option for simplicity, but it can be easily extended for the American case. With the same argument as (3.37), the value of the American option must satisfy the PDE

$$\min(\mathcal{G}[u], u(t, S, V) - g(S)) = 0 \quad (4.56)$$

### 4.3.2 PINN formulation

We can fit the variational equation (4.56) into the PINN framework established in Section 2.3 which will serve as the PDE loss (2.39), and all that is left is to set up the loss function and the relevant boundary conditions. It will be very similar to the Black-Scholes case, but as we have an additional dimension  $V$  for the variance, we will need additional boundary conditions to ensure regularity in the variance dimension as well. These boundary conditions will be inspired by those given for the European call option in [2], but we will adapt them in our study, as they exhibit different behaviour for American put options.

1.  $t \rightarrow T$ : the value at expiry is just the intrinsic value, or the payoff of the put option:

$$v(T, S, V) = \max(0, K - S) \quad (4.57)$$

which will form the terminal loss (2.40) for the PINN.

2.  $S \rightarrow 0$ : the option value is the strike when the stock is worthless, because we can exercise immediately (for European options we would require a discounting term):

$$v(t, 0, V) = K \quad (4.58)$$

3.  $S \rightarrow \infty$ : the value of a very far out of the money option is 0, as it is very unlikely for it to expire in of the money:

$$v(t, \infty, V) = 0 \quad (4.59)$$

4.  $V \rightarrow 0$ : volatility being zero is a slightly more complicated case, as this would only be true instantaneously due to the mean-reverting nature of the variance process, and the fact the variance process is guaranteed to be positive by the Feller condition. The original paper by Heston [2] proposes a collapsed first-order PDE, but this turns out to be a bit overkill for the one-dimensional case, as it doesn't yield better results. We will take this case to be that the asset price process becomes approximately deterministic, a simplification of small volatilities, and hence the value of the option can be taken to be the payoff

$$v(t, S, 0) = \max(0, K - S) \quad (4.60)$$

5.  $V \rightarrow \infty$ : when the volatility is very high, we will take it that it is almost sure that the product of assets would be zero at some point before expiry, so the value of the option is the strike:

$$v(t, S, \infty) = K \quad (4.61)$$

which is the asymptotic condition.

In practice, as we clearly cannot represent infinity, we will settle for large constants for the asset price  $S_{\inf}$  and variance  $V_{\inf}$ . The minimum for both of these will be taken to be zero. Thus our domain is

$$\Omega = (0, T) \times (0, S_{\max}) \times (0, V_{\max}) \quad (4.62)$$

The boundary conditions (2)-(5) together will form the boundary loss (2.41) for the PINN. By choosing relevant weights for each of these loss components (we could additionally set different values for different individual boundary components), we can ensure that the neural network learns the correct behaviour at the boundaries, which is crucial for the accuracy of the solution.

## 4.4 Multi-asset PINN

With the neural network established, we are now ready to move on to pricing options on multiple assets. As with the Black-Scholes case, we will price the put on the product of multiple assets because we can find a one-dimensional reduction. Here, we will consider a simplified model.

Let us have  $n$  assets with price processes  $S = (S^1, \dots, S^n)^\top$  which are correlated with each other with the  $n \times n$  correlation matrix  $\Sigma$ . Suppose further we have a common variance process



$V$  which is correlated with each of the assets. Then this process is governed by the stochastic differential equations:

$$dS = r \text{diag}(S) dt + \sqrt{V} \text{diag}(S) \text{diag}(\sigma) L dW^1 \quad (4.63)$$

$$dV = \kappa(\theta - V) dt + \bar{\sigma} \sqrt{V} dW^2 \quad (4.64)$$

where we have the vector of iid Brownian motions  $W^1 = (W^{1,1}, \dots, W^{1,n})^\top$ , volatility vector  $\sigma = (\sigma_1, \dots, \sigma_n)^\top$  and the lower-triangular Cholesky decomposition of the correlation matrix  $L$  such that  $\Sigma = LL^\top$  where its diagonal has non-negative entries. This is guaranteed to exist as  $\Sigma$  is positive semi-definite (more precisely, if  $\Sigma \succ 0$  then  $L$  exists and is unique, and if  $\Sigma \succeq 0$  then  $L$  exists but may not be unique).

For simplicity, in this model we will assume equicorrelation between the assets: that is  $\Sigma_{ij} = \varrho$  for  $i \neq j$  and  $\Sigma_{ii} = 1$ . We will require  $\varrho \in [-1, 1]$ , and be such that  $\Sigma \succeq 0$  to be a valid correlation matrix. We will also assume that each independent driving factor is correlated with the variance process:  $dW^{1,j} dW^2 = \rho_j dt$ , with  $\rho = (\rho_1, \dots, \rho_n)^\top$ , the cross-correlation vectors between assets and the variance process. Note that this means the asset  $S^i$  satisfies

$$dS^i = S^i \left( r dt + \sqrt{V} \sigma_i \sum_{j=1}^n L_{ij} dW^{1,j} \right) \quad (4.65)$$

To guarantee positiveness of the variance process, we will enforce the Feller condition on the parameters

$$2\kappa\theta \geq \bar{\sigma}^2 \quad (4.66)$$

We make these assumptions to allow a one-dimensional reduction for the Heston model, which would provide us a comparable benchmark for our results with the tree method. This would in turn give us confidence for our approximated solutions, should we wish to relax assumptions or to change the payoff.

#### 4.4.1 One-dimensional reduction

Since we would like to find an equivalent formulation of the product of assets  $X := \prod_{i=1}^n S^i$ , it would be easier to work in the log domain. That is, find the governing SDE for the log process  $Y = \log X = \sum_{i=1}^n \log S^i$  and then exponentiate with Itô's lemma to recover the desired result.

Let  $Y^i = \log S^i$ . Then by Itô's lemma, we have

$$dY^i = \frac{1}{S^i} dS^i - \frac{1}{2S^{i2}} d\langle S^i \rangle_t \quad (4.67)$$

where  $S^{i2}$  is the square of the price of the  $i$ -th asset. Here, the quadratic variation term  $d\langle S^i \rangle_t$  is given by

$$d\langle S^i \rangle_t = S^{i2} V \sigma_i^2 \sum_{j=1}^n \sum_{k=1}^n L_{ij} L_{ik} d\langle W^{1,j}, W^{1,k} \rangle_t = S^{i2} V \sigma_i^2 \sum_{j=1}^n L_{ij}^2 dt \quad (4.68)$$

and we have  $\sum_{j=1}^n L_{ij}^2 = \Sigma_{ii} = 1$ , thus giving us  $d\langle S^i \rangle_t = S^{i2} V \sigma_i^2 dt$ . Therefore, we have

$$dY^i = r dt + \sqrt{V} \sigma_i \sum_{j=1}^n L_{ij} dW^{1,j} - \frac{1}{2} V \sigma_i^2 dt \quad (4.69)$$

We can collect terms and sum over  $i$  to get (recalling  $Y = \sum_{i=1}^n Y^i$ ):

$$dY = \left( nr - \frac{1}{2} V \sum_{i=1}^n \sigma_i^2 \right) dt + \sqrt{V} \sum_{j=1}^n \sum_{i=1}^n \sigma_i L_{ij} dW^{1,j} \quad (4.70)$$

For ease of notation, let  $A = \sum_{i=1}^n \sigma_i^2$  and  $b_j = \sum_{i=1}^n \sigma_i L_{ij}$ . Then we have

$$dY = \left( nr - \frac{1}{2} V A \right) dt + \sqrt{V} \sum_{j=1}^n b_j dW^{1,j} \quad (4.71)$$

Now  $Y$  has variance

$$d\langle Y \rangle_t = V \sum_{j=1}^n b_j^2 dt =: V B dt \quad (4.72)$$

where we let  $B := \sum_{j=1}^n b_j^2$  for convenience, we can rewrite  $dY$  with a new Brownian motion:

$$dY = \left( nr - \frac{1}{2} V A \right) dt + \sqrt{V B} dW^X, \quad dW^X = \frac{\sum_{j=1}^n b_j dW^{1,j}}{\sqrt{B}} \quad (4.73)$$

Note here that  $B$  can be expressed directly in terms of  $\Sigma$ , independently of the choice of the Cholesky factor  $L$

$$B = \sum_{j=1}^n b_j^2 = \|\sigma^\top L\|^2 = \sigma^\top L L^\top \sigma = \sigma^\top \Sigma \sigma \quad (4.74)$$

which means we have

$$dW^X dW^2 = \frac{\sum_{j=1}^n b_j \rho_j}{\sqrt{B}} dt \quad (4.75)$$

so we can derive an explicit formula for the scalar correlation between  $W^X$  and  $W^2$ :

$$\rho^X = \frac{\sum_{j=1}^n b_j \rho_j}{\sqrt{B}} = \frac{\sigma^\top L \rho}{\sqrt{\sigma^\top \Sigma \sigma}} \quad (4.76)$$

Using our equicorrelation structure  $\Sigma_{ij} = \varrho$  for  $i \neq j$  and  $\Sigma_{ii} = 1$ , we have the nice simplification for  $B$ :

$$B = \sigma^\top \Sigma \sigma = (1 - \varrho) A + \varrho \left( \sum_{i=1}^n \sigma_i \right)^2 \quad (4.77)$$

Now, we have everything we need to recover  $X = e^Y$ . Using Itô's lemma again, we have

$$dX = X dY + \frac{1}{2} X d\langle Y \rangle_t \quad (4.78)$$

$$= X \left( nr + \frac{1}{2} V (B - A) \right) dt + X \sqrt{V B} dW^X \quad (4.79)$$

This is close to the Heston model, with the diffusion  $\tilde{V} = V B$  which gives us

$$d\tilde{V} = B dV \quad (4.80)$$

$$= \kappa(B\theta - \tilde{V}) dt + \bar{\sigma} \sqrt{B} \sqrt{\tilde{V}} dW^2 \quad (4.81)$$

which is a new CIR process  $\tilde{V}$  with parameters  $\tilde{\kappa} = \kappa$ ,  $\tilde{\theta} = B\theta$  and  $\tilde{\sigma} = \bar{\sigma} \sqrt{B}$  (as  $B$  is constant). The product of assets nearly gives a direct reduction to a one-dimensional Heston process, but

the drift of  $X$  contains  $V$  which is a stochastic process itself. Thus for our benchmark, we will set  $\varrho = 0$  which means  $B = A$  and the drift term becomes deterministic. The last thing to check is the Feller condition for  $\tilde{V}$ , which is satisfied as

$$2\tilde{\kappa}\tilde{\theta} \geq \tilde{\sigma}^2 \iff 2\kappa B\theta \geq \bar{\sigma}^2 B \iff 2\kappa\theta \geq \bar{\sigma}^2 \quad (4.82)$$

As such, we are able to recover a one-dimensional Heston model for the product of assets  $X$  with variance  $\tilde{V}$  and drift  $nr$ .

#### 4.4.2 Variational equation

Recall for a process  $X_t \in \mathbb{R}^n$  solving  $dX_t = b(X_t)dt + \sigma(X_t)dW_t$ , the infinitesimal generator  $\mathcal{L}$  can be derived with Itô's formula [12] on  $f(x)$  to be

$$\mathcal{L}[f](x) = \sum_{i=1}^n (b(x))_i \frac{\partial f}{\partial x_i} + \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \left( \sigma(x) \sigma(x)^\top \right)_{ij} \frac{\partial^2 f}{\partial x_i \partial x_j} \quad (4.83)$$

In our case, we have a function  $f = f(S^1, \dots, S^n, V)$ , and we can derive the infinitesimal generator much like the 1D case, but this time we need to sum over our different asset and variance processes. The drift vector is easy to see by looking at the SDEs, which is

$$b(S, V) = \begin{pmatrix} rS^1 \\ \vdots \\ rS^n \\ \kappa(\theta - V) \end{pmatrix} \quad (4.84)$$

We need to be more careful with the variance matrix, as our Brownian motions are correlated, and we need to write a variance matrix in terms of iid Brownian motions. Since  $W^{1,1}, \dots, W^{1,n}$  are already independent of each other, we only need to handle the correlation  $d\langle W^{1,j}, W^2 \rangle_t = \rho_j dt$ . Consider

$$dW^2 = \sum_{j=1}^n \rho_j dW^{1,j} + \sqrt{1 - |\rho|^2} dB^{n+1} \quad (4.85)$$

where we have  $|\rho|^2 = \sum_{j=1}^n \rho_j^2$  and where  $B^{n+1}$  is an 1-dimensional Brownian motion, independent of all  $W^{1,j}$ . Clearly we need  $|\rho|^2 \leq 1$  for this to be well-defined. Note this is quite a strong condition as it means the variance process cannot be too strongly correlated with all the assets at the same time, so care must be taken to ensure this holds before proceeding in practice.

We can easily verify that this satisfies the quadratic variation

$$d\langle W^{1,j}, W^2 \rangle_t = \rho_j dt, \quad d\langle W^2, W^2 \rangle_t = dt \quad (4.86)$$

as needed. This gives us the needed  $n + 1$  independent driving Brownian motions: we have  $(B^1, \dots, B^{n+1})$ , with  $B^i = W^{1,i}$  for  $i = 1, \dots, n$ .

We now rewrite our SDEs in the multi-asset Heston model in terms of these independent Brownian motions. The assets  $S^i$  are not driven by  $B^{n+1}$  so we simply have

$$dS^i = rS^i dt + S^i \sqrt{V} \sigma_i \sum_{j=1}^n L_{ij} dB^j + 0 \cdot dB^{n+1} \quad (4.87)$$

and then we can substitute the decomposition of  $dW^2$  into the variance process to get

$$dV = \kappa(\theta - V)dt + \bar{\sigma}\sqrt{V} \sum_{j=1}^n \rho_j dB^j + \bar{\sigma}\sqrt{V} \sqrt{1 - |\rho|^2} dB^{n+1} \quad (4.88)$$

We can now read off the  $(n+1) \times (n+1)$  variance matrix

$$\sigma(S, V) = \begin{pmatrix} \sqrt{V} \text{diag}(S) \text{diag}(\sigma) L & 0_n \\ \bar{\sigma}\sqrt{V} \rho^\top & \bar{\sigma}\sqrt{V} \sqrt{1 - |\rho|^2} \end{pmatrix} \quad (4.89)$$

which is written in block form:  $\sqrt{V} \text{diag}(S) \text{diag}(\sigma) L$  is the  $n \times n$  block for the assets,  $0_n$  is the  $n$ -dimensional zero vector for the lack of contribution of the variance Brownian motion in the assets (due to independence),  $\bar{\sigma}\sqrt{V} \rho^\top$  is the  $1 \times n$  block for the contribution of the asset Brownian motions to the variance process, and  $\bar{\sigma}\sqrt{V} \sqrt{1 - |\rho|^2}$  is the  $1 \times 1$  block for the contribution of the variance Brownian motion to the variance process. We can then compute the covariance matrix  $\tilde{\Sigma}$  (to differentiate it from  $\Sigma$  for the asset correlations) as

$$\tilde{\Sigma} = \sigma(S, V) \sigma(S, V)^\top = \begin{pmatrix} V \text{diag}(S) \text{diag}(\sigma) \Sigma \text{diag}(\sigma) \text{diag}(S) & V \bar{\sigma} \text{diag}(S) \text{diag}(\sigma) L \rho \\ V \bar{\sigma} (L \rho)^\top \text{diag}(\sigma) \text{diag}(S) & \bar{\sigma}^2 V \end{pmatrix} \quad (4.90)$$

Therefore, by plugging it into the formula (4.83), our generator can be expressed to be

$$\begin{aligned} \mathcal{L}[f] = & \sum_{i=1}^n r S^i \frac{\partial f}{\partial S^i} + \kappa(\theta - V) \frac{\partial f}{\partial V} + \frac{1}{2} V \sum_{i=1}^n \sum_{j=1}^n \sigma_i \sigma_j \Sigma_{ij} S^i S^j \frac{\partial^2 f}{\partial S^i \partial S^j} \\ & + \frac{1}{2} \bar{\sigma}^2 V \frac{\partial^2 f}{\partial V^2} + \sum_{i=1}^n \bar{\sigma} \sigma_i (L \rho)_i V S^i \frac{\partial^2 f}{\partial S^i \partial V} \end{aligned} \quad (4.91)$$

For any payoff  $g(\cdot)$ , we have, as before, the variational equation for an European option being

$$\mathcal{G}[f] := r f - \frac{\partial f}{\partial t} - \mathcal{L}[f] = 0 \quad (4.92)$$

and for an American option, we can introduce the early stopping criterion and the variation equation becomes

$$\min(\mathcal{G}[f], f - g) = 0 \quad (4.93)$$

#### 4.4.3 PINN formulation

This will be very similar to the one-dimensional case, but we will make a few modifications to account for the additional dimensions. The PDE loss will be the same as before, but the boundary conditions will include:

1.  $t \rightarrow T$ : the value at expiry is just the intrinsic value which is the payoff:

$$u(T, S^1, \dots, S^n, V) = \max \left( 0, K - \prod_{k=1}^n S^k \right) \quad (4.94)$$

which serves as our terminal loss (2.40) for the PINN.

2.  $S \rightarrow 0$ : if one of the assets reaches zero, then the product of assets is consequently zero and we should exercise immediately. Thus the value is the strike. To implement this, let  $\tilde{S}$  be  $S$  with one of the assets modified manually to 0. Then we have

$$v(t, \tilde{S}, V) = K \quad (4.95)$$

3.  $S \rightarrow \infty$ : if one of the assets has a very high price, then the product of assets would also be very high meaning the option would expire worthless out of the money. To implement this, let  $\bar{S}$  be  $S$  with one of the assets modified manually to  $S_{\text{inf}}$ , a large constant. Then the boundary condition is

$$v(t, \bar{S}, V) = 0 \quad (4.96)$$

4.  $V \rightarrow 0$ : the simplification in the one-dimensional case doesn't work well in practice for the multi-asset case, so we will consider the degenerate PDE as proposed in the original paper by Heston [2], where we have the complimentary condition

$$\min \left( rv - \frac{\partial v}{\partial t} - r \sum_{i=1}^n S^i \frac{\partial v}{\partial S^i} - \kappa \theta \frac{\partial v}{\partial V}, v - g \right) = 0 \quad (4.97)$$

where we evaluate  $v$  at  $(t, S, 0)$ .

5.  $V \rightarrow \infty$ : with infinite volatility, we could enforce the asymptotic condition as in the one-dimensional case, where

$$v(t, S, V_{\text{inf}}) = K \quad (4.98)$$

for some large constant  $V_{\text{inf}}$ . However, in practice, the PINN struggles to learn this boundary condition, and hence we will use a softer version of this constraint, taking into account the first order:


$$\frac{\partial v}{\partial V}(t, S, V_{\text{max}}) = 0 \quad (4.99)$$

using the fact that the value of the option flattens as volatility increases, since our practical implementation of  $V_{\text{max}}$  is not large enough to be in the asymptotic region, and using  $V_{\text{inf}}$  creates a great gap between the practical domain and the asymptotic region, decreasing the accuracy of the solution in the practical domain.

## 4.5 Numerical Results

### 4.5.1 Interpolating tree method

As mentioned before, we choose the resolution  $n = 100, m_v = 300, m_z = 600$  which allows a faithful approximation of closed form solution while keeping runtime low. Figure 4.3 are the results after using the parameters  $K = 1, T = 1, r = 0.05, \kappa = 2, \theta = 0.04, \sigma = 0.3, \rho = -0.7$ . We produce these plots by varying across different values of  $S_0$  (for delta),  $V_0$  (for vega) and time elapsed (for term structure), and all else kept constant.




**Figure 4.3** Call option prices by the Heston Tree

We can see that the Heston tree does indeed replicate the closed form solution very closely to very small error. This gives us confidence in our implementation, and with an easy modification we can compute the American option prices which we will use as a benchmark.

To compute the American option price with the tree model, all we need to do is to introduce the stopping condition at every node: we evaluate the intrinsic value of the option at each node and set the option value to be the maximum of the intrinsic value and the discounted continuation value.


In Figure 4.4 we can see that the European option price indeed bounds the American prices from below, as we would expect. We have confidence that our American prices are accurate since our binomial tree faithfully approximates the closed-form solution for European options. We will now use this as a benchmark for neural network approximations.



**Figure 4.4** American vs European option prices for the put


#### 4.5.2 One-dimensional PINN solution

We are now ready to train the neural network to price our option. We choose to use a learning rate of  $1 \times 10^{-3}$  with the Adam optimiser and a scheduler with step size 2000 and gamma 0.7 and the Sigmoid activation function. By using parameters  $K = 1, T = 1, r = 0.1, \kappa = 2, \theta = 0.04, \sigma = 0.3, \rho = -0.7$ , we obtain the plots as shown in Figure 4.5 for the American put.



**Figure 4.5** American put option prices by the PINN and the tree across  $S$  and  $V$  parameters at time  $t = 0$

We can see that the neural network indeed reproduces the solution approximated by the binomial tree faithfully, even though it tends to perform less well at the money especially at lower volatilities. We can take slices of our three dimensional heatmap on the triple  $(t, S, V)$  to get a better understanding of the performance, by fixing two of the variables at the default  $(t = 0, S = 1.0, V = 0.1)$  and plotting the option price against the other variable. This is shown in Figure 4.6.




**Figure 4.6** PINN vs tree prices varying across  $S$ ,  $V$  and  $t$  parameters

We can see that the PINN does indeed capture well the general shape and behaviour of the option price, especially for varying across  $S$  and  $V$  but it struggles a bit more with the term-structure (varying  $t$ ). That said, it is still able to capture the theta decay with its monotonicity, which is a good sign. Overall, the results are very promising for the one-dimensional case, with the PINN capable of capturing the option price across the whole input domain with good accuracy.

### 4.5.3 Multi-asset PINN solution

The multi-asset case is more difficult to train, as we have a higher dimensional input space and requires a greater batch size. We choose 10000 to be the batch size, with the same learning rate and step size 1500 and gamma 0.9 for the scheduler, with a L-BFGS fine-tuning step. The training takes several hours to complete and hence manual tuning and iteration of the model is slow, and I will only show the results of one instance.

As discussed before, we will set our parameters satisfying the conditions required for a one-dimensional reduction (cross-asset correlations equal to zero). The left figure of Figure 4.7 shows the heatmap of the option prices across the  $S^1$ - $S^2$  plane with the time fixed at  $t = 0$  and initial variance fixed at  $V_0 = 0.04$ .



**Figure 4.7** Heatmap of the difference between PINN and tree prices

We can see that the fit is good in the money and out of the money but struggles much more at the money (when  $S^1 \times S^2 = K$ , indicated with the black dotted line), which is expected as the option price is more sensitive, and we are now in a four-dimensional input space and hence accuracy is more difficult to achieve. We will zoom in on the at the money region by fixing  $t = 0$  and  $S^1 = S^2 = \sqrt{K}$  and plotting the option price across  $V$ , which is shown in the middle figure of Figure 4.7.

We can see that the weakness of the PINN is along the strike price with  $S^1 = S^2 = \sqrt{K}$ , and

the accuracy is better for lower volatilities. Overall, the results are okay, given that we are sacrificing some accuracy in order to be able to price along the whole input domain, but we will expect the PINN to struggle more with the estimating free boundaries correctly since there is a bias in the prices already.


Finally, we would like to check how the PINN captures the price relationship between different initial variances and time to maturity, which is shown in the right figure of Figure 4.7 where we plot the option price across  $t$  and  $V$  with  $S^1 = S^2 = \sqrt{K}$ . We can see that this relationship is captured less well: the PINN prefers to give roughly the same price across this cross-section, which is a side-effect of our loss function construction. Since the PDE loss is averaged across the domain, the network settles for giving the average price across this cross-section as it is easier to satisfy the PDE in this way, and the boundary conditions are not strong enough to force the network to learn the correct relationship across this cross-section. To remedy this, we should give more attention to how the weights of the loss components are set, which requires a lot of tuning and experimentation, which is less feasible in this case with each model taking several hours to train, but is something that could be further explored with more computational resources.

Even though the accuracy of the option price for the multi-asset case is not perfect, there is still some value in estimating the free boundary for the multi-asset case, which we will do in the next section.

#### 4.5.4 Free boundary

##### One-dimensional case

We will first present the free boundary for the one-dimensional case. As this is a boundary on three dimensions, we will present it as a cross section, where one variable will be fixed at their default value ( $t = 0, S = 1.0, V = 0.04$ ) and the other two will be plotted against each other. This would be cleaner and more interpretable than a colour-coded 3D plot. We will omit the  $t$ - $S$  plane as it will look very similar to the free boundary for the Black-Scholes case 3.15, and we will also omit the  $t$ - $V$  plane as it is not very informative, with continuation probabilities close to 1 for the whole region near the money at initial time regardless of the initial variance. This leaves us the free boundary on the  $S$ - $V$  plane, shown in Figure 4.8.



**Figure 4.8** Continuation probabilities for the American put option on the  $S$ - $V$  plane


We can see that the continuation region is reasonably well approximated with both methods, and we additionally observe some noise near the boundary which is due to the Monte Carlo estimation of the continuation value. Unfortunately, the network struggles to capture the curvature characteristic of the free boundary, which is unsurprising as it doesn't price the option with perfect accuracy in all areas. That said, the results are promising against the benchmark,



even though we have a patch of the continuation probabilities that are closer to one further in the money (at around  $(S, V) = (0.4, 0.3)$ ), which is due to the bias in the option price in the corresponding region when we had our price comparison heatmaps in Figure 4.5.

### Multi-asset case

We will first present the free boundary on the  $S^1$ - $S^2$  plane, where we fix  $V = V_0$  and  $t = 0$ , as shown in Figure 4.9. This should be somewhat analogous to the Black-Scholes multi-asset free boundary as observed in Figure 3.17.




**Figure 4.9** Continuation probabilities for the American put option on the  $S^1$ - $S^2$  plane

We can see that the continuation region is well approximated, with the shape correctly captured and the continuation probabilities resembling that of the tree method. This is promising: the PINN is able to capture the free boundary well with respect to the asset price.

We now present the free boundary on the  $S$ - $t$  plane in Figure 4.10, where we fix  $S^1 = S^2 = \sqrt{K}$  and  $V = V_0$ . This should be somewhat analogous to the free boundary for the one-dimensional case as observed in Figure 3.18. We first notice that the Heston tree free boundary is no longer smooth: this is likely due to the fact we are using an interpolation method which has some degree of error. The PINN also noticeably struggles a bit more to capture the curvature of the free boundary.

These shortcomings are likely linked to the limitations of the model, which has settled for an average price prediction across the time cross section, as an efficient and easy way of minimising the loss. To remedy this, we would need to do further experimentation with different training conditions. Even through changing the activation function between Sigmoid, Tanh and ReLU, trying different weights for the loss components, and even adding dropout layers, the results seem to be somewhat similar, so further improvements to the free boundary estimation would require a more fundamental change to the model, whether it be increasing its capacity or the sampling algorithm.



**Figure 4.10** Continuation probabilities for the American put option on the  $S$ - $t$  plane

## Chapter 5

# Conclusion

In this project, we have priced American put options on both the Black-Scholes and Heston models with various techniques, building benchmarks in the process to evaluate the validity of our results. The project mainly focuses on the use of variational methods to solve the variational inequality associated with the American option pricing problem, with analysis on the numerical results for their strengths and shortcomings. This project builds upon pre-existing techniques such as the tree and the PINN method, and adapts them to suit the American option pricing problem, while also providing a comprehensive mathematical justification for these numerical methods. We go beyond the existing literature by applying the variational methods into the problem of finding the free boundary where we develop a novel soft classification scheme, as well as constructing a solid foundation in our benchmarking process as a sanity check for our results.

We can see that the variational methods are very versatile: it is simply a matter of changing the underlying PDE and modifying the existing boundary conditions to apply the same method to a different model. Compare this with the tree method we used for the benchmark: the addition of a stochastic process for the variance in the Heston model requires a disproportionate amount of extra effort and care to construct a recombining model. The variational methods are also more efficient than the tree method: once trained, the neural network can be evaluated at any point in the domain very efficiently as it is only a forward pass, while the tree method requires the construction of the entire tree.

The versatility of the variational methods is underpinned by the exploration of the behaviour of the free boundary of the put on two assets after changing the payoffs. Even though we do not have a valid benchmark (a tree method would be difficult to construct, and a closed-form solution is of course not available), we can see that the free boundary behaves as expected. Granted, our model sacrifices some accuracy as we scale up the complexity and the number of dimensions, but the results are still reasonable, also taking into consideration hardware constraints.

A drawback with the variational method is that they are not as accurate as the tree method, since they rely on approximating the solution with only the governing PDE and boundary conditions to guide the training. We can observe systematic bias in the results in different parts of the domain space, leading to less desirable results in the free boundary approximation with the Gaussian approximation, where the unbiasedness assumption is violated. Additionally, the variational methods are not as interpretable as the tree method, which has a clear structure in its discretisation of the underlying stochastic process and its pricing architecture is more transparent with its backwards induction. As a result, the variational methods are much more difficult to debug and can actually be relatively unstable to train, as tuning various hyperparameters and changing the activation function can lead to very different results. Things brings into ques-

tion the feasibility of this method in the practical world, where the lack of interpretability and stability can be a problem for risk management, as we do not, of course, have any benchmarks to evaluate the results against.

For future work, it would be interesting to explore yet more complex models: we have shown that the variational method is versatile enough to be applied to any model with any payoff, so we could introduce either a different volatility process or add a stochastic interest rate. Additionally, the work on the multi-dimensional pricing problem could be further refined, where we take advantage of a GPU to train a larger and more capable neural network across a longer training time, which would hopefully lead to better results. As lattice methods suffer from the curse of dimensionality, we could consider using the Least-Squares Monte Carlo method as a benchmark, which is more efficient in higher dimensions. Finally, the work on the free boundary approximation could be further refined, where we work on better soft classification models and more complex Monte-Carlo schemes to retrieve the continuation value, which would hopefully allow us to observe some interesting shapes in complex payoffs in higher dimensions.

# Bibliography

- [1] F. Black and M. Scholes. The pricing of options and corporate liabilities. *Journal of political economy*, 81(3):637–654, 1973.
- [2] S. L. Heston. A closed-form solution for options with stochastic volatility with applications to bond and currency options. *The review of financial studies*, 6(2):327–343, 1993.
- [3] J. C. Cox, S. A. Ross, and M. Rubinstein. Option pricing: A simplified approach. *Journal of financial Economics*, 7(3):229–263, 1979.
- [4] M. Vellekoop and H. Nieuwenhuis. A tree-based method to price american options in the heston model. *The Journal of Computational Finance*, 13(1):1–21, 2009.
- [5] P. Glasserman. *Monte Carlo methods in financial engineering*. Springer, 2004.
- [6] M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational physics*, 378:686–707, 2019.
- [7] Y. Bai, T. Chaolu, and S. Bilige. The application of improved physics-informed neural network (ipinn) method in finance. *Nonlinear Dynamics*, 107(4):3655–3667, 2022.
- [8] Y. Zhao and H. Zheng. Fractional-boundary-regularized deep galerkin method for variational inequalities in mixed optimal stopping and control. *arXiv preprint arXiv:2505.19309*, 2025.
- [9] I. Karatzas and S. Shreve. *Brownian motion and stochastic calculus*. springer, 2014.
- [10] I. Karatzas, S. E. Shreve, I. Karatzas, and S. E. Shreve. *Methods of mathematical finance*, volume 39. Springer, 1998.
- [11] N. Wiener. Differential-space. *Journal of Mathematics and Physics*, 2(1-4):131–174, 1923.
- [12] H. Pham. *Continuous-time stochastic control and optimization with financial applications*, volume 61. Springer Science & Business Media, 2009.
- [13] Ömer Deniz Akyıldız. Stochastic simulation (lecture notes), 2025.
- [14] K. Hornik, M. Stinchcombe, and H. White. Multilayer feedforward networks are universal approximators. *Neural networks*, 2(5):359–366, 1989.
- [15] Z. Hu, K. Shukla, G. E. Karniadakis, and K. Kawaguchi. Tackling the curse of dimensionality with physics-informed neural networks. *Neural Networks*, 176:106369, 2024.
- [16] A. Dhiman and Y. Hu. Physics informed neural network for option pricing. *arXiv preprint arXiv:2312.06711*, 2023.

- [17] S. Wang, S. Sankaran, H. Wang, and P. Perdikaris. An expert’s guide to training physics-informed neural networks. *arXiv preprint arXiv:2308.08468*, 2023.
- [18] S. J. Reddi, S. Kale, and S. Kumar. On the convergence of adam and beyond. *arXiv preprint arXiv:1904.09237*, 2019.
- [19] J. Nocedal and S. J. Wright. *Numerical optimization*. Springer, 2006.
- [20] S. Markidis. The old and the new: Can physics-informed deep-learning replace traditional linear solvers? *Frontiers in big Data*, 4:669097, 2021.
- [21] D. C. Liu and J. Nocedal. On the limited memory bfgs method for large scale optimization. *Mathematical programming*, 45(1):503–528, 1989.
- [22] S. E. Shreve et al. *Stochastic calculus for finance II: Continuous-time models*, volume 11. Springer, 2004.
- [23] H. Zheng. Fundamentals of options pricing (lecture notes), 2025.
- [24] S.-S. Jo et al. Variational inequality for perpetual american option price and convergence to the solution of the difference equation. *arXiv preprint arXiv:1903.05189*, 2019.
- [25] R. A. Adams and J. J. Fournier. *Sobolev spaces*, volume 140. Elsevier, 2003.
- [26] E. Di Nezza, G. Palatucci, and E. Valdinoci. Hitchhiker’s guide to the fractional sobolev spaces. *Bulletin des sciences mathématiques*, 136(5):521–573, 2012.
- [27] Y. Zang, G. Bao, X. Ye, and H. Zhou. Weak adversarial networks for high-dimensional partial differential equations. *Journal of Computational Physics*, 411:109409, 2020.
- [28] Y. Zhao and H. Zheng. Neural network convergence for variational inequalities. *arXiv preprint arXiv:2509.26535*, 2025.
- [29] J. L. Lions and E. Magenes. *Non-homogeneous boundary value problems and applications: Vol. 1*, volume 1. Springer Science & Business Media, 2012.
- [30] J. Detemple. *American-style derivatives: Valuation and computation*. Chapman and Hall/CRC, 2005.
- [31] J. Gatheral. *The volatility surface: a practitioner’s guide*. John Wiley & Sons, 2011.
- [32] D. Duffie, J. Pan, and K. Singleton. Transform analysis and asset pricing for affine jump-diffusions. *Econometrica*, 68(6):1343–1376, 2000.
- [33] H. Albrecher, P. Mayer, W. Schoutens, and J. Tistaert. The little heston trap. *Wilmott*, (1): 83–92, 2007.
- [34] J. Gil-Pelaez. Note on the inversion theorem. *Biometrika*, 38(3-4):481–482, 1951.
- [35] W. H. Press. *Numerical recipes 3rd edition: The art of scientific computing*. Cambridge university press, 2007.
- [36] P. Bressloff. Introduction to stochastic differential equations and stochastic processes (lecture notes), 2025.